

Student Guide

awsacademy.instructure.com/courses/2401/modules/items/278437

aws

Account

Dashboard

Courses

Calendar

Inbox

History

Help

ACFv2EN-2401 > Modules > Module 5 - Networking and Content Delivery > Student Guide

Home

Announcements

Modules

Discussions

Grades

100-ACCLFO-20-EN-M05SG

< > 1 of 79 + - <->

Close

AWS Academy Cloud Foundations

Module 05 Student Guide

Version 2.0.6

100-ACCLFO-20-EN-SG

< Previous

Next >

The screenshot shows the AWS Academy Student Guide interface. The left sidebar contains navigation links: Home, Account, Dashboard, Courses, Calendar, Inbox, History, and Help. The main content area displays the title slide for 'Module 5: Networking and Content Delivery'. The slide has a blue background with the AWS Academy logo in the bottom right corner. Below the title, it says 'Welcome to Module 5: Networking and Content Delivery' and 'This module covers three fundamental Amazon Web Services (AWS) for networking and content delivery: Amazon Virtual Private Cloud (Amazon VPC), Amazon Route 53, and Amazon CloudFront.' The slide is part of a sequence of 79 items, with the current item being 4 of 79. Navigation buttons for 'Previous' and 'Next' are visible at the bottom.

The screenshot shows the 'Module overview' slide for 'Module 5: Networking and Content Delivery'. The slide is divided into three columns: Topics, Activities, and Demo. The Topics column lists: Networking basics, Amazon VPC, VPC networking, VPC security, Amazon Route 53, and Amazon CloudFront. The Activities column lists: Label a network diagram and Design a basic VPC architecture. The Demo column lists: VPC demonstration. Below these columns, there is a 'Lab' section with the task 'Build your VPC and launch a web server' and a 'Knowledge check' checkbox which is checked. The slide is part of a sequence of 79 items, with the current item being 5 of 79. Navigation buttons for 'Previous' and 'Next' are visible at the bottom.

This module addresses the following topics:

- Networking basics
- Amazon Virtual Private Cloud (Amazon VPC)
-

VPC networking

-

VPC security

-

Amazon Route 53

-

Amazon CloudFront

This module includes some activities that challenge you to label a network diagram and design a basic VPC architecture.

You will watch a recorded demonstration to learn how to use the VPC Wizard to create a VPC with public and private subnets.

You then get a chance to apply what you have learned in a hands-on lab where you use the VPC Wizard to build a VPC and launch a web server. Finally, you will be asked to complete a knowledge check that test your understanding of key concepts that are covered in this module.

After completing this module, you should be able to

:

Recognize the basics of networking

-

Describe virtual networking in the cloud with Amazon VPC

-

Label a network diagram

-

Design a basic VPC architecture

-

Indicate the steps to build a VPC

-

Identify security groups

-

Create your own VPC and add additional components to it to produce a customized network

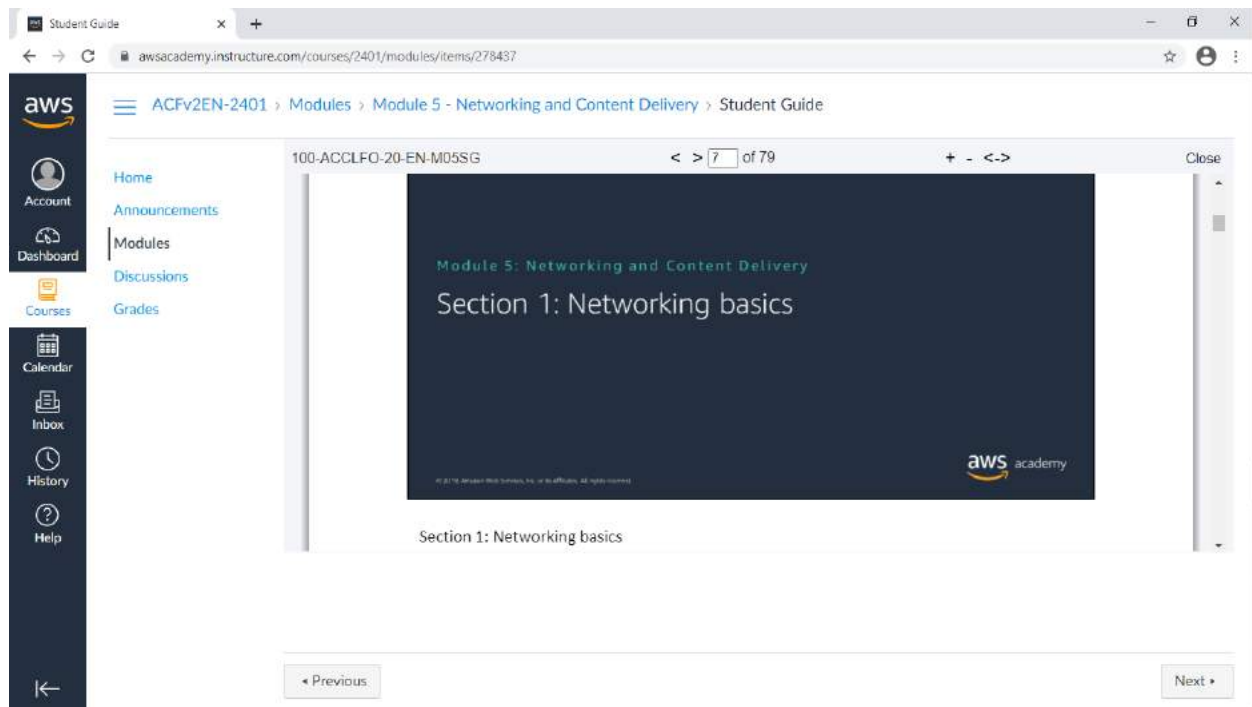
-

Identify the fundamentals of Amazon Route 53

-

Recognize the benefits of Amazon CloudFront

6



Section 1: Networking basics

In this section, you will review a few basic networking concepts that provide the necessary foundation to your understanding of the AWS networking service, Amazon Virtual Private Cloud (Amazon VPC).

Student Guide

awsacademy.instructure.com/courses/2401/modules/items/278437

ACFv2EN-2401 > Modules > Module 5 - Networking and Content Delivery > Student Guide

Account

Dashboard

Courses

Calendar

Inbox

History

Help

Home

Announcements

Modules

Discussions

Grades

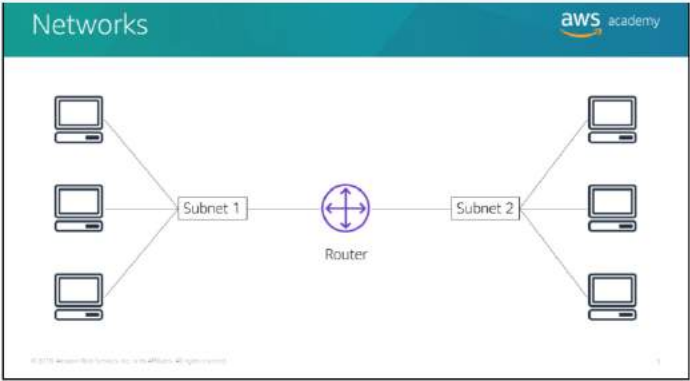
100-ACCLFO-20-EN-M05SG

< > 8 of 79

+ - <->

Close

Networks



```
graph LR; S1[Subnet 1] --- R((Router)); R --- S2[Subnet 2]; S1 --- C1[Computer]; S1 --- C2[Computer]; S1 --- C3[Computer]; S2 --- C4[Computer]; S2 --- C5[Computer]; S2 --- C6[Computer];
```

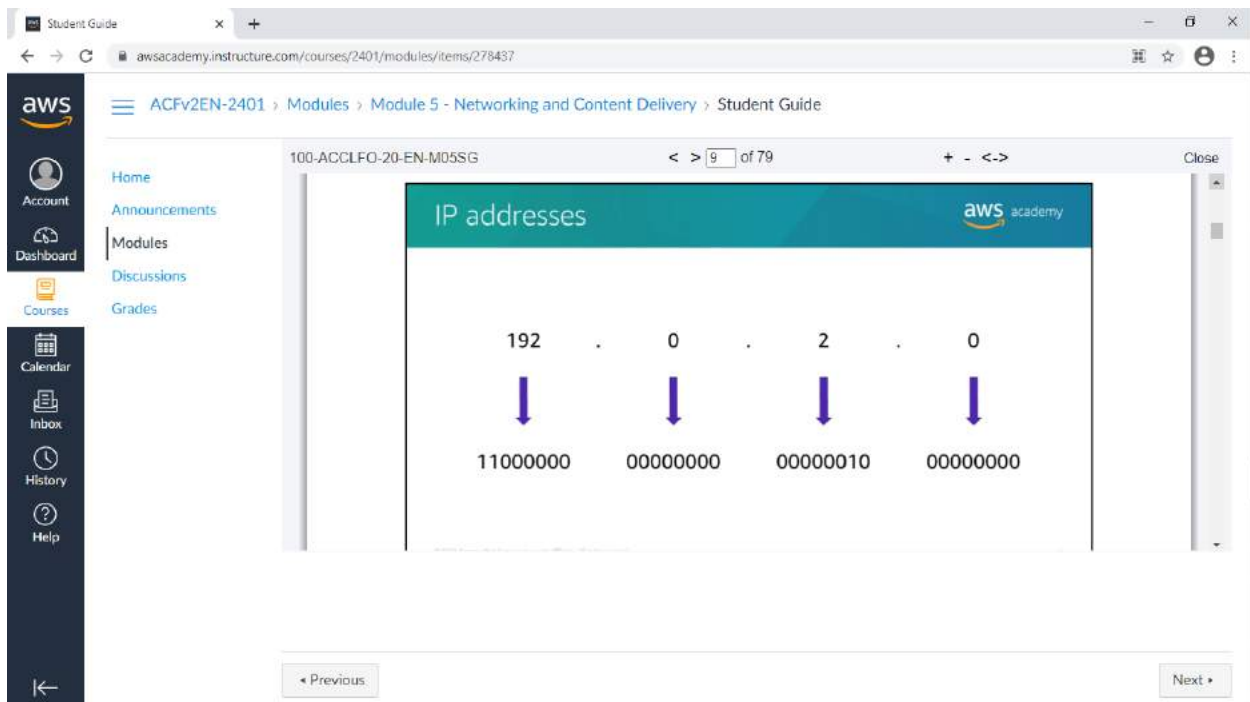
© 2020 Amazon Web Services, Inc. or its affiliates. All rights reserved.

< Previous

Next >

A computer network is two or more client machines that are connected together to share resources. A network can be logically partitioned into subnets

. Networking requires a networking device (such as a router or switch) to connect all the clients together and enable communication between them.



Each client machine in a network has a unique Internet Protocol (IP) address that identifies it.

An IP address is a numerical label in decimal format. Machines convert that decimal number to a binary format.

In this example, the IP address is 192.0.2.0. Each of the four dot (.)-separated numbers of the IP address represents 8 bits in octal number format

. That means each of the four numbers can be anything from 0 to 255. The combined total of the four numbers for an IP address is 32 bits in binary format.

Student Guide

awsacademy.instructure.com/courses/2401/modules/items/278437

aws

Account

Dashboard

Courses

Calendar

Inbox

History

Help

ACFv2EN-2401 > Modules > Module 5 - Networking and Content Delivery > Student Guide

Home

Announcements

Modules

Discussions

Grades

100-ACCLFO-20-EN-M05SG

< > 10 of 79

+ - <->

Close

IPv4 and IPv6 addresses

aws academy

IPv4 (32-bit) address: 192.0.2.0

IPv6 (128-bit) address: 2600:1f18:22ba:8c00:ba86:a05e:a5ba:00ff

© 2019 Amazon Web Services, Inc. or its affiliates. All rights reserved.

← Previous

Next →

A 32-bit IP address is called an IPv4 address.

IPv6 addresses, which are 128 bits, are also available. IPv6 addresses can accommodate more user devices.

An IPv6 address is composed of eight groups of four letters and numbers that are separated

by colons (:). In this example, the IPv6 address is 2600:1f18:22ba:8c00:ba86:a05e:a5ba:00FF.

Each of the eight colon-separated groups of the IPv6 address represents 16 bits in hexadecimal number format. That means each of the eight groups can be anything from 0 to FFFF. The combined total of the eight groups for an IPv6 address is 128 bits in binary format.

The screenshot shows a web browser window displaying the AWS Academy interface. The browser's address bar shows the URL `awsacademy.instructure.com/courses/2401/modules/items/278437`. The page title is "Student Guide". The left sidebar contains navigation links: Home, Account, Announcements, Modules, Discussions, Grades, Calendar, Inbox, History, and Help. The main content area displays a slide titled "Classless Inter-Domain Routing (CIDR)" with the AWS Academy logo. The slide shows the IP address `192.0.2.0/24` and explains its components: the network identifier (routing prefix) `192.0.2` and the host identifier `0`. Below the IP address, the binary representation of each octet is shown: `11000000` (Fixed), `00000000` (Fixed), `00000010` (Fixed), and `00000000 to 11111111` (Flexible). A callout indicates that the `24` in the slash notation tells how many bits are fixed. The slide is part of a presentation, with navigation buttons for "Previous" and "Next" at the bottom.

A common method to describe networks is Classless Inter-Domain Routing (CIDR). The CIDR address is expressed as follows:

-

An IP address (which is the first address of the network)

- t

Next, a slash character (/)

-

Finally, a number that tells you how many bits of the routing prefix must be fixed or allocated for the network identifier

The bits that are not fixed are allowed to change.

CIDR is a way to express a group of IP

addresses that are consecutive to each other.

In this example, the CIDR address is 192.0.2.0/24.

The last number (24) tells you that the first 24 bits must be fixed. The last 8 bits are flexible, which means that 2⁸ (or 256) IP addresses are available for the network, which range from 192.0.2.0 to 192.0.2.255.

The fourth decimal

digit is allowed to change from 0 to 255.

If the CIDR was 192.0.2.0/16, the last number (16)

tells you that the first 16 bits must be

fixed. The last 16 bits are flexible, which means that 2¹⁶

(or 65,536) IP addresses are available

for the network, ranging from 192.0.0.0 to 192.0.255.255. The third and fourth decimal digits

can each change from 0 to 255.

There are two special cases:

-

Fixed IP addresses, in which every bit is fixed, represent a single IP address (for example, 192.0.2.0/32). This type of address is helpful when you want to set up a firewall rule and give access to a specific host.

-

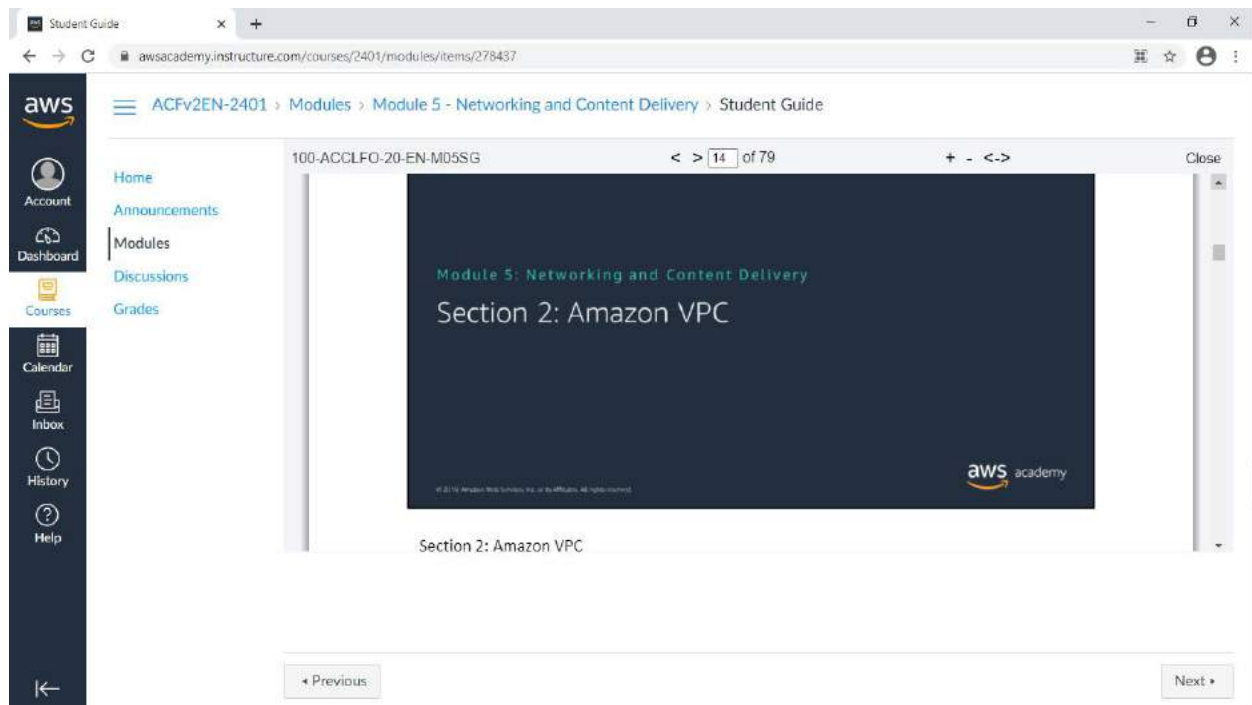
The internet, in which every bit is flexible, is represented as 0.0.0.0/0

The screenshot shows a web browser window with the URL `awsacademy.instructure.com/courses/2401/modules/items/278437`. The page is titled "Student Guide" and is part of the "ACFv2EN-2401" course, specifically "Module 5 - Networking and Content Delivery". The left sidebar contains navigation links: Home, Announcements, Modules, Discussions, Grades, Account, Dashboard, Courses, Calendar, Inbox, History, and Help. The main content area displays a slide titled "Open Systems Interconnection (OSI) model" from AWS Academy. The slide features a table with the following data:

| Layer | Number | Function | Protocol/Address |
|--------------|--------|---|--------------------------|
| Application | 7 | Means for an application to access a computer network | HTTP(S), FTP, DHCP, LDAP |
| Presentation | 6 | <ul style="list-style-type: none">Ensures that the application layer can read the dataEncryption | ASCII, ICA |
| Session | 5 | Enables orderly exchange of data | NetBIOS, RPC |
| Transport | 4 | Provides protocols to support host-to-host communication | TCP, UDP |
| Network | 3 | Routing and packet forwarding (routers) | IP |
| Data link | 2 | Transfer data in the same LAN network (hubs and switches) | MAC |
| Physical | 1 | Transmission and reception of raw bitstreams over a physical medium | Signals (1s and 0s) |

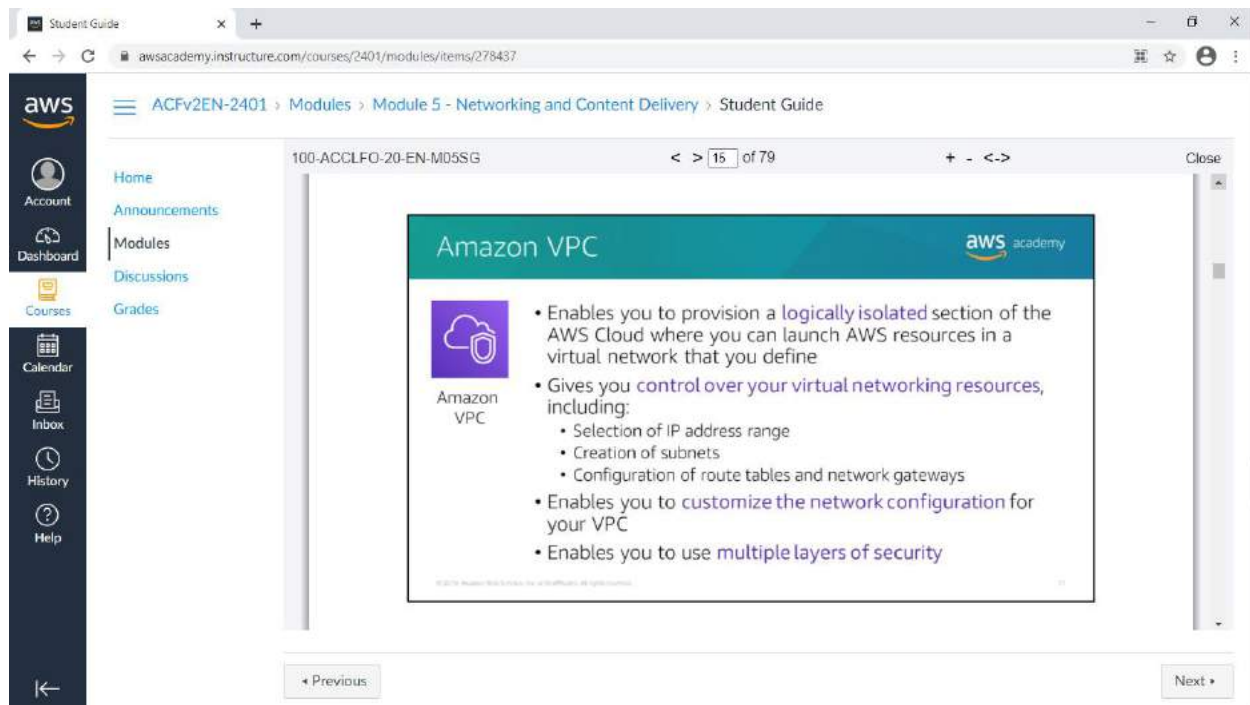
At the bottom of the slide, there is a small copyright notice: "© 2019 Amazon Web Services, Inc. or its affiliates. All rights reserved." The slide is part of a sequence of 79 items, with the current item being 13. Navigation buttons for "Previous" and "Next" are visible at the bottom of the slide frame.

The Open Systems Interconnection (OSI) model is a conceptual model that is used to explain how data travels over a network. It consists of seven layers and shows the common protocols and addresses that are used to send data at each layer. For example, hubs and switches work at layer 2 (the data link layer). Routers work at layer 3 (the network layer). The OSI model can also be used to understand how communication takes place in a virtual private cloud (VPC), which you will learn about in the next section.



Section 2: Amazon VPC

Many of the concepts of an on-premises network apply to a cloud-based network, but much of the complexity of setting up a network has been abstracted without sacrificing control, security, and usability. In this section, you learn about Amazon VPC and the fundamental components of a VPC.



Amazon Virtual Private Cloud (Amazon VPC) is a service that lets you provision a logically isolated section of the AWS Cloud (called a virtual private cloud, or VPC) where you can launch your AWS resources.

Amazon VPC gives you control over your virtual networking resources, including the selection of your own IP address range, the creation of subnets, and the configuration of route tables and network gateways. You can use both IPv4 and IPv6 in your VPC for secure access to resources and applications.

You can also customize the network configuration for your VPC. For example, you can create a public subnet for your web servers that can access the public internet. You can place your backend systems (such as databases or applications servers) in a private subnet with no public internet access.

Finally, you can use multiple layers of security, including security groups and network access control lists (network ACLs), to help control access to Amazon Elastic Compute Cloud

(Amazon EC2) instances in each subnet.

Student Guide

awsacademy.instructure.com/courses/2401/modules/items/278437

ACFv2EN-2401 > Modules > Module 5 - Networking and Content Delivery > Student Guide

Account

Dashboard

Courses

Calendar

Inbox

History

Help

Home

Announcements

Modules

Discussions

Grades

100-ACCLFO-20-EN-M05SG < > 16 of 79 + - <-> Close

VPCs and subnets

- VPCs:
 - Logically isolated from other VPCs
 - Dedicated to your AWS account
 - Belong to a single AWS Region and can span multiple Availability Zones
- Subnets:
 - Range of IP addresses that divide a VPC
 - Belong to a single Availability Zone
 - Classified as public or private

The diagram illustrates the AWS Cloud architecture. It shows a large box labeled 'AWS Cloud' containing a 'Region' box. Inside the 'Region' box, there are two 'Availability Zone' boxes labeled 'Availability Zone 1' and 'Availability Zone 2'. A 'VPC' box spans across both 'Availability Zone' boxes. Inside each 'Availability Zone' box, there is a 'Subnet' box.

< Previous

Next >

Amazon VPC enables you to provision virtual private clouds (VPCs). A VPC is a virtual network that is logically isolated from other virtual networks in the AWS Cloud. A VPC is dedicated to your account. VPCs belong to a single AWS Region and can span multiple Availability Zones.

After you create a VPC, you can divide it into one or more subnets. A subnet is a range of IP addresses in a VPC. Subnets belong to a single Availability Zone. You can create subnets in different Availability Zones for high availability.

Subnets are generally classified as public or private. Public subnets have direct access to the internet, but private subnets do not.

Student Guide

awsacademy.instructure.com/courses/2401/modules/items/278437

ACFv2EN-2401 > Modules > Module 5 - Networking and Content Delivery > Student Guide

Account

Dashboard

Courses

Calendar

Inbox

History

Help

Home

Announcements

Modules

Discussions

Grades

100-ACCLFO-20-EN-M05SG < > 17 of 79 + - <-> Close

IP addressing

aws academy

- When you create a VPC, you assign it to an IPv4 CIDR block (range of private IPv4 addresses).
- You cannot change the address range after you create the VPC.
- The largest IPv4 CIDR block size is /16.
- The smallest IPv4 CIDR block size is /28.
- IPv6 is also supported (with a different block size limit).
- CIDR blocks of subnets cannot overlap.

VPC

x.x.x.x/16 or 65,536 addresses (max)
to
x.x.x.x/28 or 16 addresses (min)

© 2019 Amazon Web Services, Inc. or its affiliates. All rights reserved.

Previous

Next

IP addresses enable resources in your VPC to communicate with each other and with resources over the internet. When you create a VPC, you assign an IPv4 CIDR block (a range of private IPv4 addresses) to it. After you create a VPC, you cannot change the address range, so it is important that you choose it carefully. The IPv4 CIDR block might be as large as /16 (which is 2^{16} or 65,536 addresses) or as small as /28 (which is 2^4 , or 16 addresses). You can optionally associate an IPv6 CIDR block with your VPC and subnets, and assign IPv6 addresses from that block to the resources in your VPC. IPv6 CIDR blocks have a different block size limit. The CIDR block of a subnet can be the same as the CIDR block for a VPC. In this case, the VPC and the subnet are the same size (a single subnet in the VPC). Also, the CIDR block of a subnet can be a subset of the CIDR block for the VPC. This structure enables the definition of multiple subnets. If you create more than one subnet in a VPC, the CIDR blocks of the subnets cannot overlap. You cannot have duplicate IP addresses in the same VPC. To learn more about IP addressing in a VPC, see the [AWS Documentation](#).

Student Guide

awsacademy.instructure.com/courses/2401/modules/items/278437

ACFv2EN-2401 > Modules > Module 5 - Networking and Content Delivery > Student Guide

Account

Dashboard

Courses

Calendar

Inbox

History

Help

Home

Announcements

Modules

Discussions

Grades

100-ACCLFO-20-EN-M05SG < > 18 of 79 + - <-> Close

Reserved IP addresses

aws academy

Example: A VPC with an IPv4 CIDR block of 10.0.0.0/16 has 65,536 total IP addresses. The VPC has four equal-sized subnets. Only 251 IP addresses are available for use by each subnet.

VPC: 10.0.0.0/16

Subnet 1 (10.0.0.0/24)

251 IP addresses

Subnet 2 (10.0.2.0/24)

251 IP addresses

Subnet 4 (10.0.1.0/24)

251 IP addresses

Subnet 3 (10.0.3.0/24)

251 IP addresses

| IP Addresses for CIDR block 10.0.0.0/24 | Reserved for |
|---|-------------------------------------|
| 10.0.0.0 | Network address |
| 10.0.0.1 | Internal communication |
| 10.0.0.2 | Domain Name System (DNS) resolution |
| 10.0.0.3 | Future use |
| 10.0.0.255 | Network broadcast address |

Previous

Next

When you create a subnet, it requires its own CIDR block. For each CIDR block that you specify, AWS reserves five IP addresses within that block, and these addresses are not available for use. AWS reserves these IP addresses for:

-

Network address

-

VPC local router (internal communications)

-

Domain Name System (DNS) resolution

-

Future use

-

Network broadcast address

For example, suppose that you create a subnet with an IPv4 CIDR block of 10.0.0.0/24 (which has 256 total IP addresses). The subnet has 256 IP addresses, but only 251 are available because five are reserved.

Student Guide

awsacademy.instructure.com/courses/2401/modules/items/278437

aws

Account

Dashboard

Courses

Calendar

Inbox

History

Help

ACFv2EN-2401

Modules

Module 5 - Networking and Content Delivery

Student Guide

Home

Announcements

Modules

Discussions

Grades

100-ACCLFO-20-EN-M05SG

< > 19 of 79

+ - <->

Close

Public IP address types

Public IPv4 address

- Manually assigned through an Elastic IP address
- Automatically assigned through the auto-assign public IP address settings at the subnet level

Elastic IP address

- Associated with an AWS account
- Can be allocated and remapped anytime
- Additional costs might apply

© 2019 Amazon Web Services, Inc. or its affiliates. All rights reserved.

79

Previous

Next

When you create a VPC, every instance in that VPC gets a private IP address automatically.

You can also request a public IP address to be assigned when you create the instance by modifying the subnet's auto-assign public IP address properties.

An Elastic IP address is a static and public IPv4 address that is designed for dynamic cloud computing. You can associate an Elastic IP address with any instance or network interface for

any VPC in your account. With an Elastic IP address

, you can mask the failure of an instance

by rapidly remapping the address to another instance in your VPC.

Associating the Elastic IP

address with the network interface has an advantage

over associating it directly with the

instance. You can move all of the attributes of the

network interface from one instance to

another in a single step.

Additional costs might apply when you use Elastic IP addresses, so it is important to release

them when you no longer need them.

To learn more about Elastic IP addresses, see

Elastic IP Addresses

in the AWS Documentation.

Student Guide


awsacademy.instructure.com/courses/2401/modules/items/278437

ACFv2EN-2401 > Modules > Module 5 - Networking and Content Delivery > Student Guide

100-ACCLFO-20-EN-M05SG < > 20 of 79 + - <-> Close

Elastic network interface

- An elastic network interface is a **virtual network interface** that you can:
 - Attach to an instance.
 - Detach from the instance, and attach to another instance to redirect network traffic.
- Its **attributes follow** when it is reattached to a new instance.
- Each instance in your VPC has a **default network interface** that is assigned a private IPv4 address from the IPv4 address range of your VPC.



© 2019 Amazon Web Services, Inc. or its affiliates. All rights reserved.

← Previous Next →

An elastic network interface is a virtual network interface that you can attach or detach from an instance in a VPC. A network interface's attributes follow it when it is reattached to another instance. When you move a network interface from one instance to another, network traffic is redirected to the new instance. Each instance in your VPC has a default network interface (the primary network interface) that is assigned a private IPv4 address from the IPv4 address range of your VPC. You cannot detach a primary network interface from an instance. You can create and attach an additional network interface to any instance in your VPC. The number of network interfaces you can attach varies by instance type. For more information about Elastic Network Interfaces, see the AWS Documentation.

The screenshot shows the AWS Academy Student Guide interface. The browser address bar displays `awsacademy.instructure.com/courses/2401/modules/items/278437`. The page title is "Route tables and routes". The left sidebar contains navigation links: Home, Announcements, Modules, Discussions, Grades, Calendar, Inbox, History, and Help. The main content area is titled "Route tables and routes" and includes a list of bullet points and a table titled "Main (Default) Route Table".

- A route table contains a set of rules (or routes) that you can configure to direct network traffic from your subnet.
- Each route specifies a destination and a target.
- By default, every route table contains a local route for communication within the VPC.
- Each subnet must be associated with a route table (at most one).

| Destination | Target |
|-------------|--------|
| 10.0.0.0/16 | local |

A line points from the text "VPC CIDR block" to the "10.0.0.0/16" entry in the table.

A route table contains a set of rules (called routes) that directs network traffic from your subnet. Each route specifies a destination and a target. The destination is the destination CIDR block where you want traffic from your subnet to go. The Target is the target that the destination traffic is sent through. By default, every route table that you create contains a local route for communication in the VPC. You can customize route tables by adding routes.

You cannot delete the local route entry that is used for internal communications.

Each subnet in your VPC must be associated with a route table. The main route table is the route table is automatically assigned to your VPC. It controls the routing for all subnets that are not explicitly associated with any other route table. A subnet can be associated with only one route table at a time, but you can associate multiple subnets with the same route table.

To learn more about route tables, see the AWS Documentation

Student Guide

awsacademy.instructure.com/courses/2401/modules/items/278437

ACFv2EN-2401 > Modules > Module 5 - Networking and Content Delivery > Student Guide

100-ACCLFO-20-EN-M05SG < > 22 of 79 + - <-> Close

Home
Announcements
Modules
Discussions
Grades

Account
Dashboard
Courses
Calendar
Inbox
History
Help

Section 2 key takeaways

aws academy

- A VPC is a logically isolated section of the AWS Cloud.
- A VPC belongs to one Region and requires a CIDR block.
- A VPC is subdivided into subnets.
- A subnet belongs to one Availability Zone and requires a CIDR block.
- Route tables control traffic for a subnet.
- Route tables have a built-in local route.
- You add additional routes to the table.
- The local route cannot be deleted.

Previous Next

Some key takeaways from this section of the module include:

- A VPC is a logically isolated section of the AWS Cloud.
- A VPC belongs to one Region and requires a CIDR block.
- A VPC is subdivided into subnets.
- A subnet belongs to one Availability Zone and requires a CIDR block.
- Route tables control traffic for a subnet.
- Route tables have a built-in local route.
- You add additional routes to the table.
- The local route cannot be deleted.

Student Guide

awsacademy.instructure.com/courses/2401/modules/items/278437

ACFv2EN-2401 > Modules > Module 5 - Networking and Content Delivery > Student Guide

Home
Announcements
Modules
Discussions
Grades

100-ACCLFO-20-EN-M05SG < > 23 of 79 + - <-> Close

Module 5: Networking and Content Delivery
Section 3: VPC networking

© 2019 Amazon Web Services, Inc. or its affiliates. All rights reserved.

Previous Next

Section 3: VPC networking

Now that you have learned about the basic components of a VPC, you can start routing traffic in interesting ways. In this section, you learn about different networking options.

Student Guide

awsacademy.instructure.com/courses/2401/modules/items/278437

ACFv2EN-2401 > Modules > Module 5 - Networking and Content Delivery > Student Guide

Home
Announcements
Modules
Discussions
Grades

100-ACCLFO-20-EN-M05SG < > 24 of 79 + - <-> Close

Internet gateway

AWS Cloud

Region

Availability Zone

VPC: 10.0.0.0/16

Public subnet: 10.0.1.0/24

Private subnet: 10.0.2.0/24

Route table

Internet gateway (igw-id)

Internet

| Destination | Target |
|-------------|--------|
| 10.0.0.0/16 | local |
| 0.0.0.0/0 | igw-id |

Previous Next

An

internet gateway

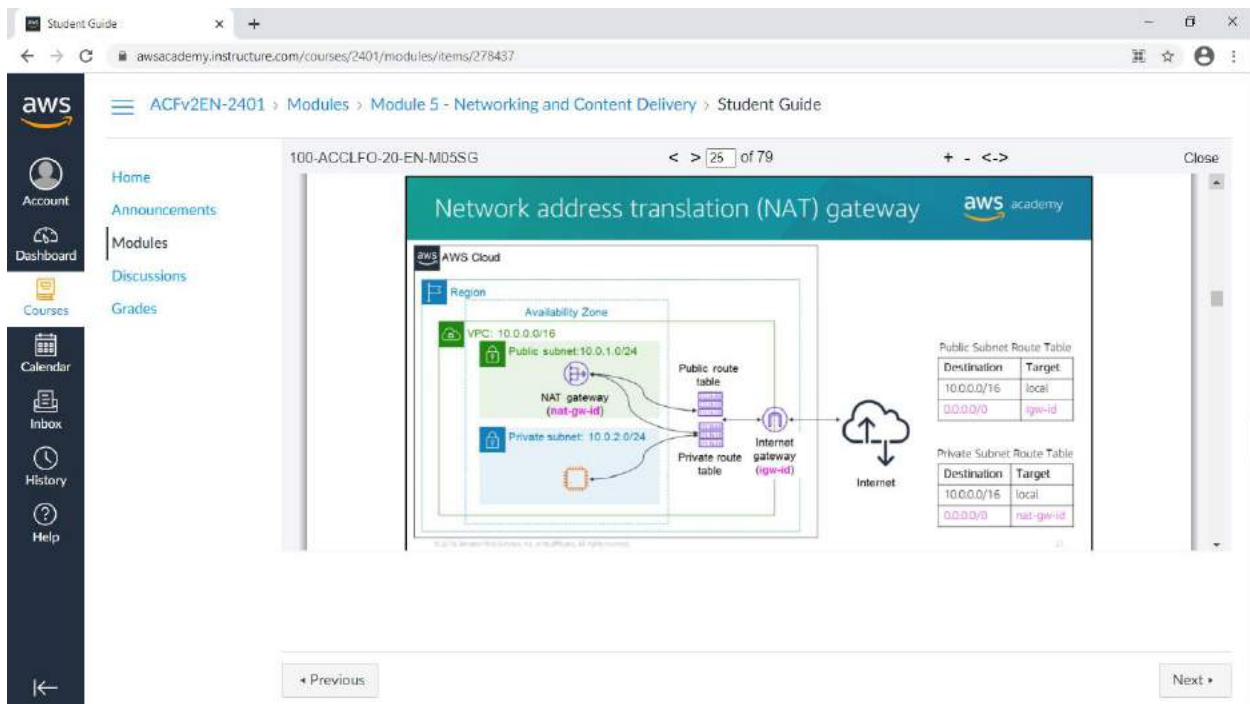
is a scalable, redundant, and highly available VPC component that allows

communication between instances in your VPC and the internet. An internet gateway serves

two purposes: to provide a target in your VPC route tables for internet-routable traffic, and to perform network address translation for instances that were assigned public IPv4 addresses.

To make a subnet public, you attach an internet gateway to your VPC and add a route to the route table to send non-local traffic through the internet gateway to the internet (0.0.0.0/0).

For more information about internet gateways, see Internet Gateways in the AWS Documentation.



A network address translation (NAT) gateway enables instances in a private subnet to connect to the internet or other AWS services, but prevents the internet from initiating a connection with those instances.

To create a NAT gateway, you must specify the public subnet in which the NAT gateway should reside. You must also specify an Elastic IP address to associate with the NAT gateway when you create it. After you create a NAT gateway, you must update the route table that is associated with one or more of your private subnets to point internet-bound traffic to the NAT gateway. Thus, instances in your private subnets can communicate with the internet.

You can also use a NAT instance in a public subnet in your VPC instead of a NAT gateway.

However, a NAT gateway is a managed NAT service that provides better availability, higher bandwidth, and less administrative effort. For common use cases, AWS recommends that you use a NAT gateway instead of a NAT instance.

See the AWS Documentation for more information about

-

NAT gateways

-

NAT instances

-

Differences between NAT gateways and NAT instances

Student Guide

awsacademy.instructure.com/courses/2401/modules/items/278437

ACFv2EN-2401 > Modules > Module 5 - Networking and Content Delivery > Student Guide

Home

Announcements

Modules

Discussions

Grades

Account

Dashboard

Courses

Calendar

Inbox

History

Help

100-ACCLFO-20-EN-M05SG

< > 26 of 79

+

-

<->

Close

VPC sharing

aws academy

aws

AWS Cloud

Region

VPC: Account A (owner)

Private subnet

Public subnet

Router

Account B (participant)

Account C (participant)

Account D (participant)

EC2 instance

EC2 instance

EC2 instance

RDS instance

NAT gateway

EC2 instance

Amazon Redshift

Internet gateway

< Previous

Next >

VPC sharing enables customers to share subnets with other AWS accounts in the same organization in AWS Organizations. VPC sharing enables multiple AWS accounts to create their application resources—such as Amazon EC2 instances, Amazon Relational Database Service (Amazon RDS) databases, Amazon Redshift clusters, and AWS Lambda functions—into shared, centrally managed VPCs. In this model, the account that owns the VPC (owner) shares one or more subnets with other accounts (participants) that belong to the same organization in AWS Organizations. After a subnet is shared, the participants can view, create, modify, and delete their application resources in the subnets that are shared with them. Participants cannot view, modify, or delete resources that belong to other participants or the VPC owner.

VPC sharing offers several benefits:

- Separation of duties – Centrally controlled VPC structure, routing, IP address allocation
- Ownership – Application owners continue to own resources, accounts, and security groups
- Security groups – VPC sharing participants can reference the security group IDs of each other
- Efficiencies – Higher density in subnets, efficient use of VPNs and AWS Direct Connect
- No hard limits – Hard limits can be avoided—for example, 50 virtual interfaces per

AWS Direct Connect connection through simplified network architecture

-

Optimized costs – Costs can be optimized through the reuse of NAT gateways, VPC interface endpoints, and intra-Availability Zone traffic. VPC sharing enables you to decouple accounts and networks. You have fewer, larger, centrally managed VPCs. Highly interconnected applications automatically benefit from this approach.

Student Guide

awsacademy.instructure.com/courses/2401/modules/items/278437

ACFv2EN-2401 > Modules > Module 5 - Networking and Content Delivery > Student Guide

100-ACCLFO-20-EN-M05SG < > 28 of 79 + - <-> Close

Home
Announcements
Modules
Discussions
Grades

Account
Dashboard
Courses
Calendar
Inbox
History
Help

VPC peering

aws academy

aws AWS Cloud

VPC A: 10.0.0.0/16

VPC B: 10.3.0.0/16

Peering connection (pcx-id)

You can connect VPCs in your own AWS account, between AWS accounts, or between AWS Regions.

Restrictions:

- IP spaces cannot overlap.
- Transitive peering is not supported.
- You can only have one peering resource between the same two VPCs.

Route Table for VPC A

| Destination | Target |
|-------------|--------|
| 10.0.0.0/16 | local |
| 10.3.0.0/16 | pcx-id |

Route Table for VPC B

| Destination | Target |
|-------------|--------|
| 10.3.0.0/16 | local |
| 10.0.0.0/16 | pcx-id |

Previous Next

A

VPC peering connection

is a networking connection between two VPCs that enables you to route traffic between them privately. Instances in either VPC can communicate with each other as if they are within the same network. You can create a VPC peering connection between your own VPCs, with a VPC in another AWS account, or with a VPC in a different AWS Region.

When you set up the peering connection, you create rules in your route table to allow the VPCs to communicate with each other through the peering resource. For example, suppose that you have two VPCs. In the route table for VPC A, you set the destination to be the IP address of VPC B and the target to be the peering resource ID. In the route table for VPC B, you set the destination to be the IP address of VPC A and the target to be the peering resource ID.

VPC peering has some restrictions:

-

IP address ranges cannot overlap.

-

Transitive peering is not supported. For example, suppose that you have three VPCs: A, B, and C. VPC A is connected to VPC B, and VPC A is connected to VPC C. However, VPC B is not connected to VPC C implicitly. To connect VPC B to VPC C, you must explicitly establish that connectivity.

-

You can only have one peering resource between the same two VPCs.

For more information about VPC peering, see [VPC Peering](#) in the AWS Documentation.

Student Guide

awsacademy.instructure.com/courses/2401/modules/items/278437

Home

Account

Dashboard

Courses

Calendar

Inbox

History

Help

ACFv2EN-2401 > Modules > Module 5 - Networking and Content Delivery > Student Guide

100-ACCLFO-20-EN-M05SG < > 29 of 79 + - < -> Close

AWS Site-to-Site VPN

The diagram illustrates an AWS Site-to-Site VPN setup. It shows a VPC with two subnets: a Public subnet (10.0.0.0/24) and a Private subnet (10.0.2.0/24). A Route table is associated with the VPC, and a Virtual gateway (vgw-1a) is connected to the Internet. A Site-to-Site VPN connection is established between the Virtual gateway and a Customer gateway. The Corporate data center is shown with the IP address 192.168.10.0/24. Two route tables are also displayed: the Public subnet route table and the Private subnet route table.

| Destination | Target |
|-------------|--------|
| 10.0.0.0/16 | local |
| 0.0.0.0/0 | igw-id |

| Destination | Target |
|-----------------|--------|
| 10.0.0.0/16 | local |
| 192.168.10.0/24 | vgw-id |

Corporate data center: 192.168.10.0/24

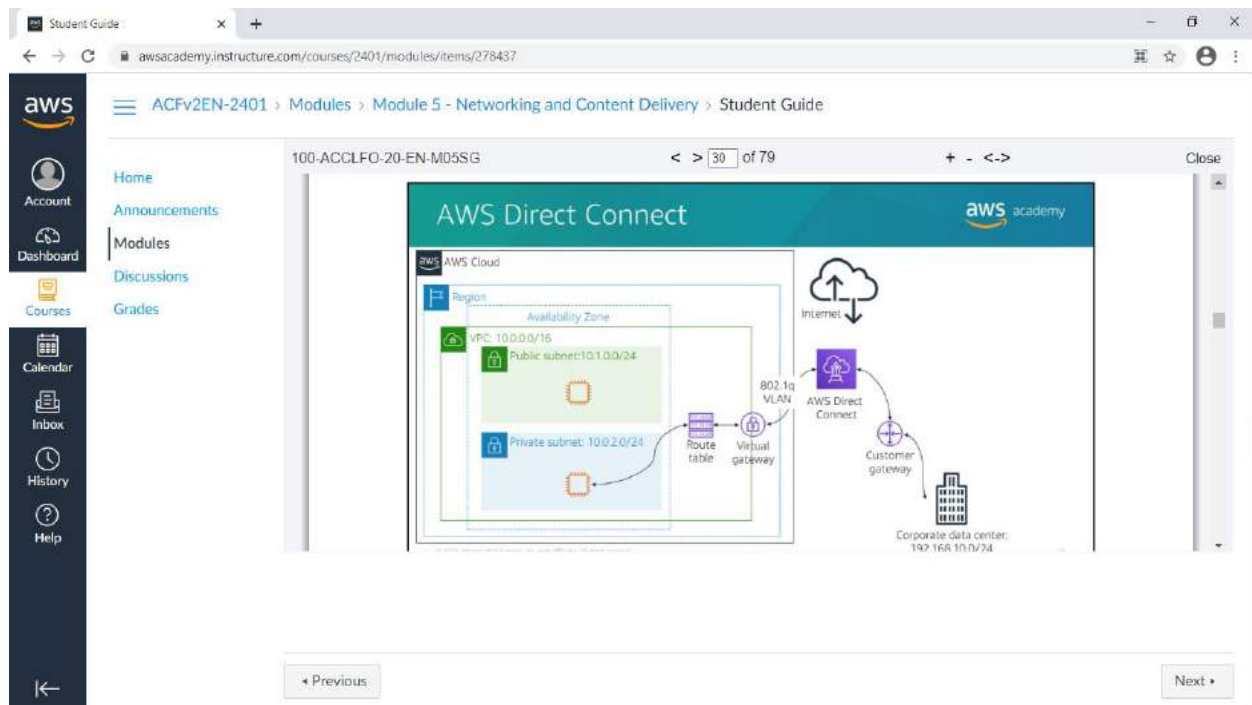
< Previous

Next >

By default, instances that you launch into a VPC cannot communicate with a remote network.

To connect your VPC to your remote network (that is, create a virtual private network or VPN connection), you:

1. Create a new virtual gateway device (called a virtual private network (VPN) gateway) and attach it to your VPC.
 2. Define the configuration of the VPN device or the customer gateway.
. The customer gateway is not a device but an AWS resource that provides information to AWS about your VPN device.
 3. Create a custom route table to point corporate data center-bound traffic to the VPN gateway. You also must update security group rules. (You will learn about security groups in the next section.)
 4. Establish an AWS Site-to-Site VPN (Site-to-Site VPN) connection to link the two systems together.
 5. Configure routing to pass traffic through the connection.
- For more information about AWS Site-to-Site VPN and other VPN connectivity options, see [VPN Connections](#) in the AWS Documentation.



One of the challenges of network communication is network performance. Performance can be negatively affected if your data center is located far away from your AWS Region. For such situations, AWS offers AWS Direct Connect, or DX. AWS Direct Connect enables you to establish a dedicated, private network connection between your network and one of the DX locations. This private connection can reduce your network costs, increase bandwidth throughput, and provide a more consistent network experience than internet-based connections. DX uses open standard 802.1q virtual local area networks (VLANs). For more information about DX, see the AWS Direct Connect product page

Student Guide

awsacademy.instructure.com/courses/2401/modules/items/278437

ACFv2EN-2401 > Modules > Module 5 - Networking and Content Delivery > Student Guide

100-ACCLFO-20-EN-M05SG < > 31 of 79 + - <-> Close

VPC endpoints

Public Subnet Route Table

| Destination | Target |
|-------------|---------|
| 10.0.0.0/16 | local |
| Amazon S3 | vpc-elb |

Two types of endpoints:

- Interface endpoints (powered by AWS PrivateLink)
- Gateway endpoints (Amazon S3 and Amazon DynamoDB)

Previous Next

A

VPC endpoint is a virtual device that enables you to privately connect your VPC to supported AWS services and VPC endpoint services that are powered by AWS PrivateLink.

Connection to these services does not require an internet gateway, NAT device, VPN connection, or AWS Direct Connect connection. Instances in your VPC do not require public IP addresses to communicate with resources in the service. Traffic between your VPC and the other service does not leave the Amazon network.

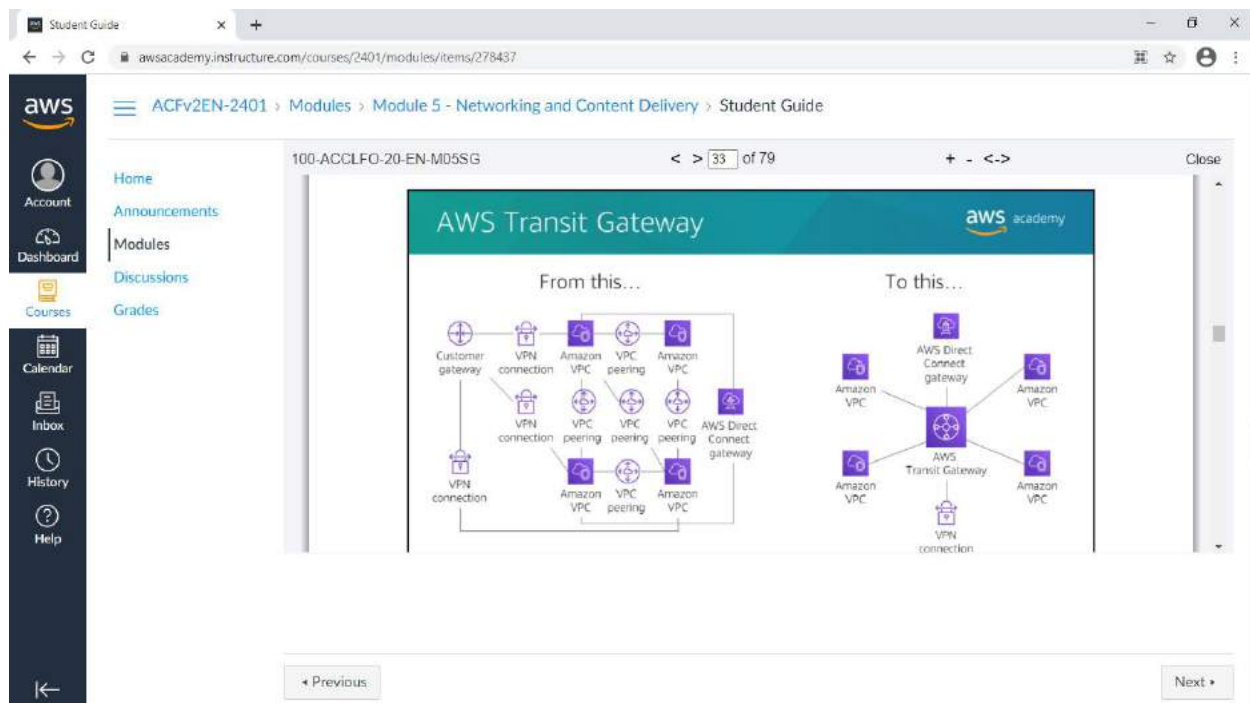
There are two types of VPC endpoints:

-

An interface VPC endpoint (interface endpoint) enables you to connect to services that are powered by AWS PrivateLink. These services include some AWS services, services that are hosted by other AWS customers and AWS Partner Network (APN) Partners in their own VPCs (referred to as endpoint services), and supported AWS Marketplace APN Partner

services. The owner of the service is the service provider, and you—as the principal who creates the interface endpoint—are the service consumer. You are charged for creating and using an interface endpoint to a service. Hourly usage rates and data processing rates apply. See the AWS Documentation for a list of supported interface endpoints and for more information about the example shown here.

- Gateway endpoints: The use of gateway endpoints incurs no additional charge. Standard charges for data transfer and resource usage apply. For more information about VPC endpoints, see VPC Endpoints in the AWS Documentation.



You can configure your VPCs in several ways, and take advantage of numerous connectivity options and gateways. These options and gateways include AWS Direct Connect (via DX gateways), NAT gateways, internet gateways, VPC peering, etc. It is not uncommon to find AWS customers with hundreds of VPCs distributed across AWS accounts and Regions to serve multiple lines of business, teams, projects, and so forth. Things get more complex when customers start to set up connectivity between their VPCs. All the connectivity options are strictly point-to-point, so the number of VPC-to-VPC connections can grow quickly. As you grow the number of workloads that run on AWS, you must be able to scale your networks across multiple accounts and VPCs to keep up with the growth. Though you can use VPC peering to connect pairs of VPCs, managing point-to-point connectivity across many VPCs without the ability to centrally manage the connectivity policies can be operationally costly and difficult. For on-premises connectivity, you must attach your VPN to each individual VPC. This solution can be time-consuming to build and difficult to manage when the number of VPCs grows into the hundreds. To solve this problem, you can use AWS Transit Gateway to simplify your networking model. With AWS Transit Gateway, you only need to create and manage a single connection from the central gateway into each VPC, on-premises data center, or

remote office across your network. A transit gateway acts as a hub that controls how traffic is routed among all the connected networks, which act like spokes. This hub-and-spoke model significantly simplifies management and reduces operational costs because each network only needs to connect to the transit gateway and not to every other network. Any new VPC is connected to the transit gateway, and is then automatically available to every other network that is connected to the transit gateway. This ease of connectivity makes it easier to scale your network as you grow.

Student Guide

awsacademy.instructure.com/courses/2401/modules/items/278437

Home

Account

Dashboard

Courses

Calendar

Inbox

History

Help

ACFv2EN-2401 > Modules > Module 5 - Networking and Content Delivery > Student Guide

100-ACCLFO-20-EN-M05SG < > 35 of 79 + - <-> Close

Activity: Label this network diagram

| Destination | Target |
|-------------|--------|
| 0.0.0.0/0 | local |

< Previous

Next >

See if you can recognize the different VPC networking components that you learned about by labeling this network diagram.

Student Guide

awsacademy.instructure.com/courses/2401/modules/items/278437

ACFv2EN-2401 > Modules > Module 5 - Networking and Content Delivery > Student Guide

100-ACCLFO-20-EN-M05SG < > 36 of 79 + - < - > Close

Activity: Solution

| Destination | Target |
|-------------|--------|
| 10.0.0.0/16 | local |
| 0.0.0.0/0 | igw-1d |

Previous Next

Now, see how well you did.

Student Guide

awsacademy.instructure.com/courses/2401/modules/items/278437

ACFv2EN-2401 > Modules > Module 5 - Networking and Content Delivery > Student Guide

100-ACCLFO-20-EN-M05SG < > 37 of 79 + - < - > Close

Recorded Amazon VPC demonstration

Set up demo
Amazon Virtual Private Cloud (VPC)

Previous Next

Now that you know how to design a VPC, watch this demonstration to learn how to use the VPC Wizard to set up a VPC with public and private subnets.

The screenshot displays the AWS Academy Student Guide interface. The browser address bar shows the URL `awsacademy.instructure.com/courses/2401/modules/items/278437`. The page title is "Student Guide". The breadcrumb navigation path is "ACFv2EN-2401 > Modules > Module 5 - Networking and Content Delivery > Student Guide".

The left sidebar contains navigation links: Home, Announcements, Modules, Discussions, Grades, Account, Dashboard, Courses, Calendar, Inbox, History, and Help.

The main content area shows a slide titled "Section 3 key takeaways" with a key icon. The slide content includes:

- There are several VPC networking options, which include:
 - Internet gateway
 - NAT gateway
 - VPC endpoint
 - VPC peering
 - VPC sharing
 - AWS Site-to-Site VPN
 - AWS Direct Connect
 - AWS Transit Gateway
- You can use the VPC Wizard to implement your design.

The slide is part of a presentation titled "100-ACCLFO-20-EN-M05SG" and is slide 38 of 79. Navigation buttons for "Previous" and "Next" are visible at the bottom.

Some key takeaways from this section of the module include:

-

There are several VPC networking options, which include:

-

Internet gateway: Connects your VPC to the internet

-

NAT gateway: Enables instances in a private subnet to connect to the internet

-

VPC endpoint: Connects your VPC to supported AWS services

-

VPC peering: Connects your VPC to other VPCs

-

VPC sharing: Allows multiple AWS accounts to create their application resources into shared, centrally-managed Amazon VPCs

-

AWS Site-to-Site VPN: Connects your VPC to remote networks

-

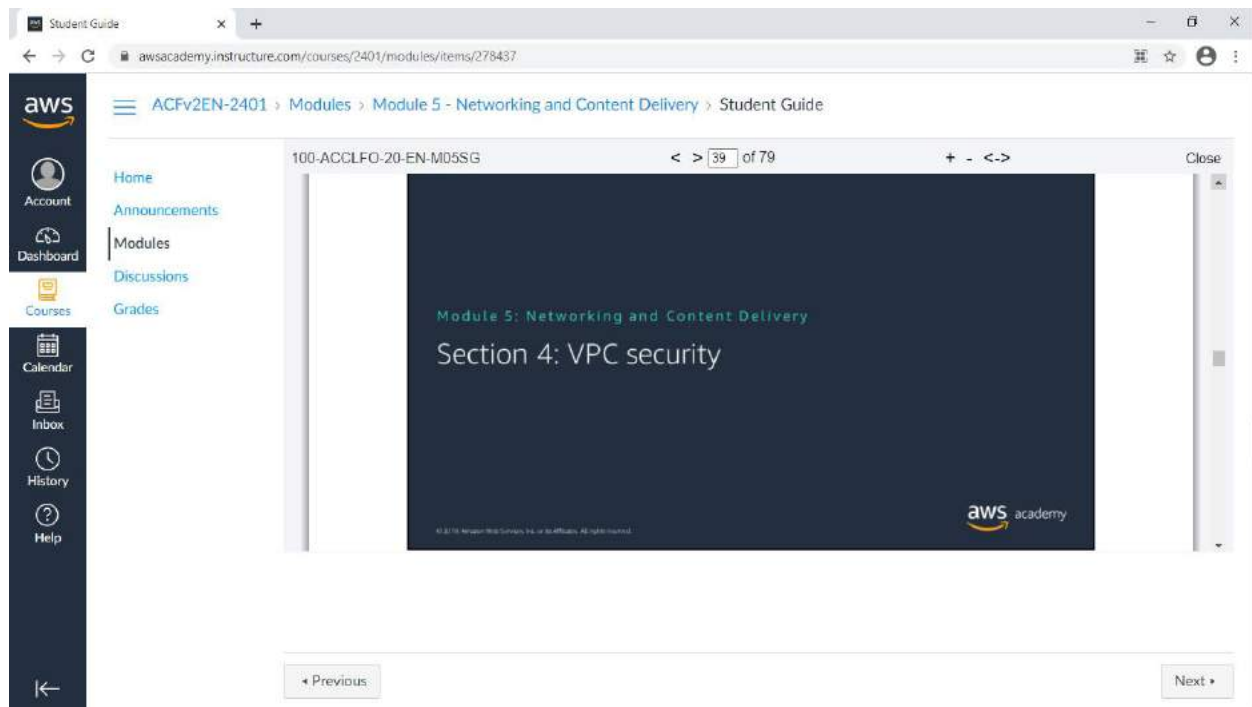
AWS Direct Connect: Connects your VPC to a remote network by using a dedicated network connection

-

AWS Transit Gateway: A hub-and-spoke connection alternative to VPC peering

-

You can use the VPC Wizard to implement your design



Section 4: VPC security

You can build security into your VPC architecture in several ways so that you have complete control over both incoming and outgoing traffic. In this section, you learn about two Amazon VPC firewall options that you can use to secure your VPC: security groups and network access control lists (network ACLs).

Student Guide

awsacademy.instructure.com/courses/2401/modules/items/278437

ACFv2EN-2401 > Modules > Module 5 - Networking and Content Delivery > Student Guide

Account

Dashboard

Courses

Calendar

Inbox

History

Help

Home

Announcements

Modules

Discussions

Grades

100-ACCLFO-20-EN-M05SG

< > 40 of 79

+ - <->

Close

Security groups

Security groups act at the instance level.

< Previous

Next >

A security group acts as a virtual firewall for your instance, and it controls inbound and outbound traffic. Security groups act at the instance level, not the subnet level. Therefore, each instance in a subnet in your VPC can be assigned to a different set of security groups. At the most basic level, a security group is a way for you to filter traffic to your instances.

Student Guide

awsacademy.instructure.com/courses/2401/modules/items/278437

ACFv2EN-2401 > Modules > Module 5 - Networking and Content Delivery > Student Guide

aws

Account

Dashboard

Courses

Calendar

Inbox

History

Help

Home

Announcements

Modules

Discussions

Grades

100-ACCLFO-20-EN-M05SG < > 41 of 79 + - <-> Close

Security groups

| Type | Protocol | Port Range | Source | Description |
|-------------|----------|------------|-------------|-------------|
| All traffic | All | All | sg-xxxxxxxx | |

| Type | Protocol | Port Range | Source | Description |
|-------------|----------|------------|-------------|-------------|
| All traffic | All | All | sg-xxxxxxxx | |

- Security groups have rules that control inbound and outbound instance traffic.
- Default security groups deny all inbound traffic and allow all outbound traffic.
- Security groups are stateful.

< Previous

Next >

Security groups have rules that control the inbound and outbound traffic. When you create a security group, it has no inbound rules. Therefore, no inbound traffic that originates from another host to your instance is allowed until you add inbound rules to the security group.

By default, a security group includes an outbound rule that allows all outbound traffic

. You can remove the rule and add outbound rules that allow specific outbound traffic only. If your security group has no outbound rules, no outbound traffic that originates from your instance is allowed.

Security groups are stateful

, which means that state information is kept even after a request is processed. Thus, if you send a request from your instance, the response traffic for that request is allowed to flow in regardless of inbound security group rules. Responses to allowed inbound traffic are allowed to flow out, regardless of outbound rules.

Student Guide

awsacademy.instructure.com/courses/2401/modules/items/278437

ACFv2EN-2401 > Modules > Module 5 - Networking and Content Delivery > Student Guide

100-ACCLFO-20-EN-M05SG < > 42 of 79 + - <-> Close

Custom security groups

| Inbound | | | | |
|---------|----------|------------|----------------|-----------------|
| Type | Protocol | Port Range | Source | Description |
| HTTP | TCP | 80 | 0.0.0.0/0 | All web traffic |
| HTTPS | TCP | 443 | 0.0.0.0/0 | All web traffic |
| SSH | TCP | 22 | 54.24.12.19/32 | Office address |

| Outbound | | | | |
|-------------|----------|------------|-----------|-------------|
| Type | Protocol | Port Range | Source | Description |
| All traffic | All | All | 0.0.0.0/0 | |
| All traffic | All | All | :::0 | |

- You can specify allow rules, but not deny rules.
- All rules are evaluated before the decision to allow traffic.

Previous Next

When you create a custom security group, you can specify allow rules, but not deny rules. All rules are evaluated before the decision to allow traffic.

42

Student Guide

awsacademy.instructure.com/courses/2401/modules/items/278437

ACFv2EN-2401 > Modules > Module 5 - Networking and Content Delivery > Student Guide

100-ACCLFO-20-EN-M05SG < > 43 of 79 + - <-> Close

Network access control lists (network ACLs)

The diagram illustrates the AWS Cloud architecture. It shows a Region containing an Availability Zone. Within the Availability Zone, there is a VPC (Virtual Private Cloud) with the IP address 10.0.0.0/16. The VPC contains two subnets: a Public subnet (10.0.0.0/24) and a Private subnet (10.0.4.0/22). Network ACLs (Access Control Lists) are shown as small blue squares with a lock icon, positioned between the subnets and the VPC. A callout box points to these Network ACLs with the text: "Network ACLs act at the subnet level."

Previous Next

A network access control list (network ACL) is an optional layer of security for your Amazon VPC. It acts as a firewall for controlling traffic in and out of one or more subnets. To add another layer of security to your VPC, you can set up network ACLs with rules that are similar to your security groups.

Each subnet in your VPC must be associated with a network ACL. If you don't explicitly

associate a subnet with a network ACL, the subnet is automatically associated with the default network ACL. You can associate a network ACL with multiple subnets; however, a subnet can be associated with only one network ACL at a time. When you associate a network ACL with a subnet, the previous association is removed.

Student Guide

awsacademy.instructure.com/courses/2401/modules/items/278437

ACFv2EN-2401 > Modules > Module 5 - Networking and Content Delivery > Student Guide

100-ACCLFO-20-EN-M05SG < > 44 of 79 + - <-> Close

Network ACLs

| Inbound | | | | | |
|---------|------------------|----------|------------|-----------|------------|
| Rule # | Type | Protocol | Port Range | Source | Allow/Deny |
| 100 | All IPv4 traffic | All | All | 0.0.0.0/0 | ALLOW |
| * | All IPv4 traffic | All | All | 0.0.0.0/0 | DENY |

| Outbound | | | | | |
|----------|------------------|----------|------------|-----------|------------|
| Rule # | Type | Protocol | Port Range | Source | Allow/Deny |
| 100 | All IPv4 traffic | All | All | 0.0.0.0/0 | ALLOW |
| * | All IPv4 traffic | All | All | 0.0.0.0/0 | DENY |

- A network ACL has separate inbound and outbound rules, and each rule can either allow or deny traffic.
- Default network ACLs allow all inbound and outbound IPv4 traffic.
- Network ACLs are stateless.

© 2019 Amazon Web Services, Inc. or its affiliates. All rights reserved.

◀ Previous Next ▶

A network ACL has separate inbound and outbound rules, and each rule can either allow or deny traffic. Your VPC automatically comes with a modifiable default network ACL. By default, it allows all inbound and outbound IPv4 traffic and, if applicable, IPv6 traffic. The table shows a default network ACL. Network ACLs are stateless, which means that no information about a request is maintained after a request is processed.

Student Guide

awsacademy.instructure.com/courses/2401/modules/items/278437

ACFv2EN-2401 > Modules > Module 5 - Networking and Content Delivery > Student Guide

Account

Dashboard

Courses

Calendar

Inbox

History

Help

Home

Announcements

Modules

Discussions

Grades

100-ACCLFO-20-EN-M05SG

< > 45 of 79

+ - <->

Close

Custom network ACLs

| Inbound | | | | | |
|---------|------------------|----------|------------|-----------|------------|
| Rule # | Type | Protocol | Port Range | Source | Allow/Deny |
| 103 | SSH | TCP | 22 | 0.0.0.0/0 | ALLOW |
| 100 | HTTPS | TCP | 443 | 0.0.0.0/0 | ALLOW |
| * | All IPv4 traffic | All | All | 0.0.0.0/0 | DENY |

| Outbound | | | | | |
|----------|------------------|----------|------------|-----------|------------|
| Rule # | Type | Protocol | Port Range | Source | Allow/Deny |
| 103 | SSH | TCP | 22 | 0.0.0.0/0 | ALLOW |
| 100 | HTTPS | TCP | 443 | 0.0.0.0/0 | ALLOW |
| * | All IPv4 traffic | All | All | 0.0.0.0/0 | DENY |

- Custom network ACLs deny all inbound and outbound traffic until you add rules.
- You can specify both allow and deny rules.
- Rules are evaluated in number order, starting with the lowest number.

< Previous

Next >

You can create a custom network ACL and associate it with a subnet. By default, each custom network ACL denies all inbound and outbound traffic until you add rules.

A network ACL contains a numbered list of rules that are evaluated in order, starting with the lowest numbered rule. The purpose is to determine whether traffic is allowed in or out of any subnet that is associated with the network ACL. The highest number that you can use for a rule is 32,766. AWS recommends that you create rules in increments (for example, increments of 10 or 100) so that you can insert new rules where you need them later.

For more information about network ACLs, see [Network ACLs in the AWS Documentation](#).

The screenshot shows the AWS Academy Student Guide interface. The browser address bar displays the URL: `awsacademy.instructure.com/courses/2401/modules/items/278437`. The page title is "100-ACCLFO-20-EN-M05SG". The left sidebar contains navigation links: Home, Account, Dashboard, Courses, Calendar, Inbox, History, and Help. The main content area displays a slide titled "Security groups versus network ACLs" with the AWS Academy logo. The slide contains a table comparing the two concepts.

| Attribute | Security Groups | Network ACLs |
|-----------------|---|--|
| Scope | Instance level | Subnet level |
| Supported Rules | Allow rules only | Allow and deny rules |
| State | Stateful (return traffic is automatically allowed, regardless of rules) | Stateless (return traffic must be explicitly allowed by rules) |
| Order of Rules | All rules are evaluated before decision to allow traffic | Rules are evaluated in number order before decision to allow traffic |

Navigation controls at the bottom of the slide include "Previous" and "Next" buttons. The slide footer includes the text "© 2020 Amazon Web Services, Inc. or its affiliates. All rights reserved." and the number "18".

Here is a summary of the differences between security groups and network ACLs:

- Security groups act at the instance level, but network ACLs act at the subnet level.
- Security groups support allow rules only, but network ACLs support both allow and deny rules.
- Security groups are stateful, but network ACLs are stateless.
- For security groups, all rules are evaluated before the decision is made to allow traffic.
For network ACLs, rules are evaluated in number order before the decision is made to allow traffic.

Student Guide

awsacademy.instructure.com/courses/2401/modules/items/278437

aws

Account

Dashboard

Courses

Calendar

Inbox

History

Help

ACFv2EN-2401 > Modules > Module 5 - Networking and Content Delivery > Student Guide

Home

Announcements

Modules

Discussions

Grades

100-ACCLFO-20-EN-M05SG

< > 47 of 79

+ - <->

Close

Activity: Design a VPC

Scenario: You have a small business with a website that is hosted on an Amazon Elastic Compute Cloud (Amazon EC2) instance. You have customer data that is stored on a backend database that you want to keep private. You want to use Amazon VPC to set up a VPC that meets the following requirements:

- Your web server and database server must be in separate subnets.
- The first address of your network must be 10.0.0.0. Each subnet must have 256 total IPv4 addresses.
- Your customers must always be able to access your web server.
- Your database server must be able to access the internet to make patch updates.
- Your architecture must be highly available and use at least one custom firewall layer.

Now, it's your turn! In this scenario, you are a small business owner with a website that is hosted on an Amazon Elastic Compute Cloud (Amazon EC2) instance. You have customer data

< Previous

Next >

Now, it's your turn! In this scenario, you are a small business owner with a website that is hosted on an Amazon Elastic Compute Cloud (Amazon EC2) instance. You have customer data that is stored on a backend database that you want to keep private.

See if you can design a VPC that meets the following requirements:

-

Your web server and database server must be in separate subnets.

-

The first address of your network must be 10.0.0.0.

Each subnet must have 256 IPv4 addresses.

-

Your customers must always be able to access your web server.

-

Your database server must be able to access the internet to make patch updates.

-

Your architecture must be highly available and use at least one custom firewall layer.

The screenshot shows the AWS Academy Student Guide interface. The browser address bar displays `awsacademy.instructure.com/courses/2401/modules/items/278437`. The left sidebar contains navigation links: Home, Account, Dashboard, Courses, Calendar, Inbox, History, and Help. The main content area shows a slide titled "Section 4 key takeaways" with a key icon and the word "Takeaway". The slide content includes the following bullet points:

- Build security into your VPC architecture:
 - Isolate subnets if possible.
 - Choose the appropriate gateway device or VPN connection for your needs.
 - Use firewalls.
- Security groups and network ACLs are firewall options that you can use to secure your VPC.

The slide is part of a presentation titled "100-ACCLFO-20-EN-M05SG" and "Module 5 - Networking and Content Delivery". The slide number is 48 of 79. Navigation buttons for "Previous" and "Next" are visible at the bottom.

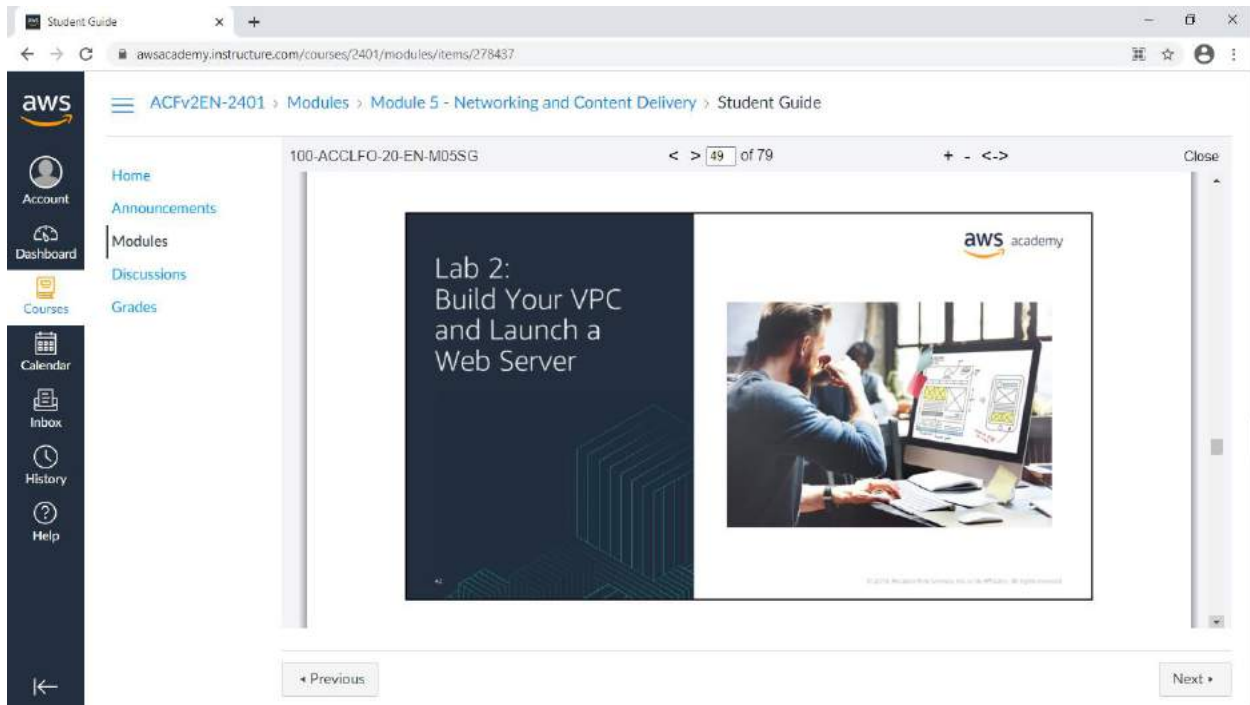
The key takeaways from this section of the module are:

-

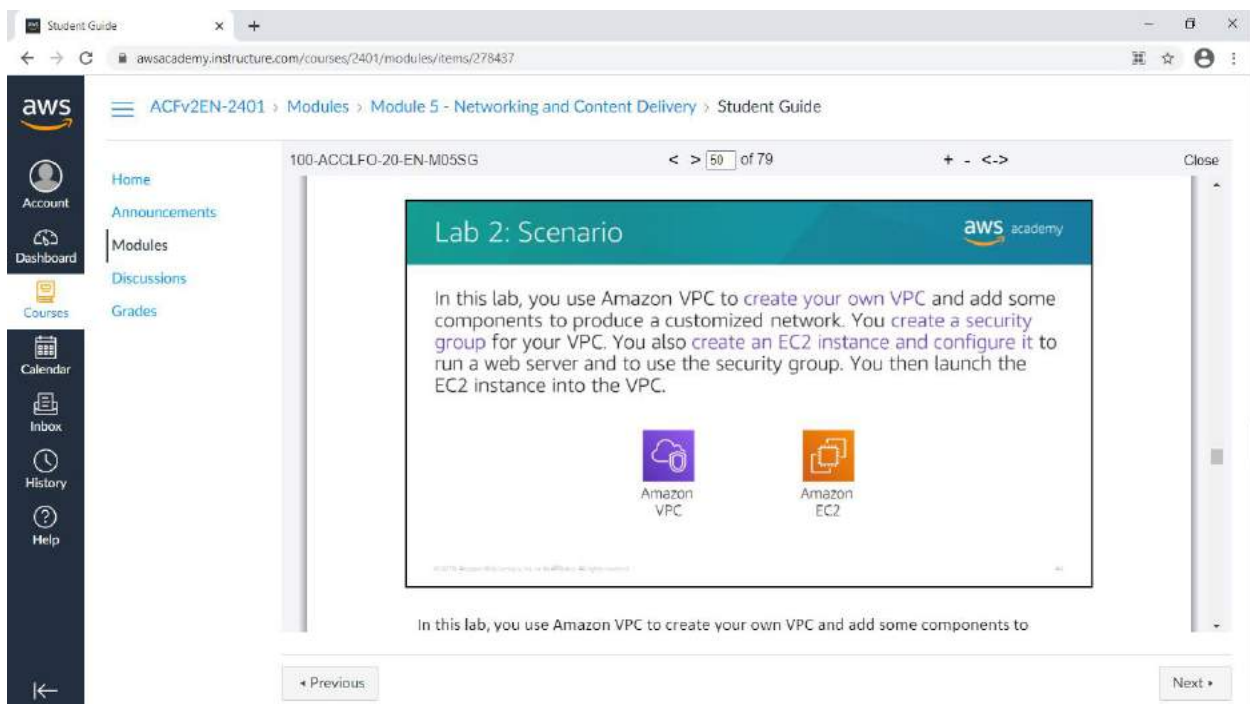
Build security into your VPC architecture.

-

Security groups and network ACLs are firewall options that you can use to secure your VPC.



You will now work on Lab 2: Build Your VPC and Launch a Web Server.
49

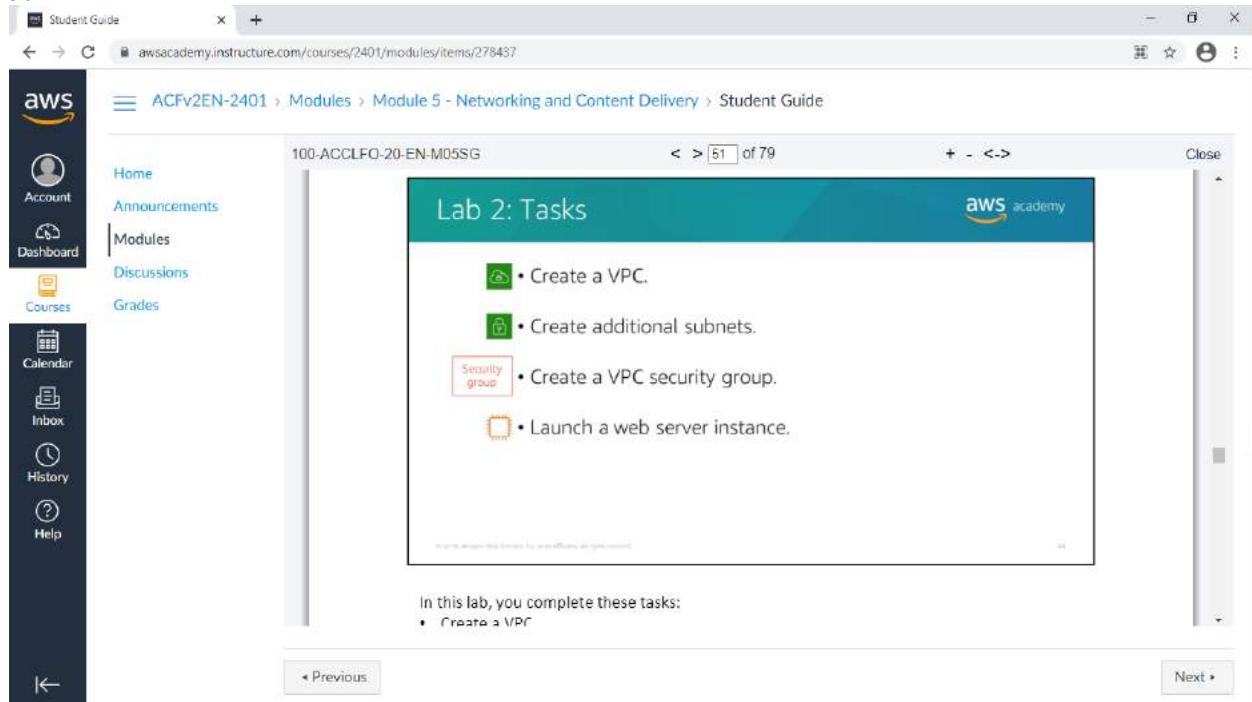


In this lab, you use Amazon VPC to create your own VPC and add some components to

produce a customized network. You also create a security group for your VPC, and then create an EC2 instance and configure it to run a web server and to use the security group.

You then launch the EC2 instance into the VPC.

50



In this lab, you complete these tasks:

- Create a VPC.
- Create additional subnets.
- Create a VPC security group.
- Launch a web server instance.

51

Student Guide

awsacademy.instructure.com/courses/2401/modules/items/278437

ACFv2EN-2401 > Modules > Module 5 - Networking and Content Delivery > Student Guide

Home

Announcements

Modules

Discussions

Grades

Account

Dashboard

Courses

Calendar

Inbox

History

Help

100-ACCLFO-20-EN-M05SG

< > 52 of 79

+ - <->

Close

Lab 2: Final product

© 2019 Amazon Web Services, Inc. or its affiliates. All rights reserved.

| Destination | Target |
|-------------|------------------|
| 10.0.0.0/16 | Local |
| 0.0.0.0/0 | Internet gateway |

| Destination | Target |
|-------------|-------------|
| 10.0.0.0/16 | Local |
| 0.0.0.0/0 | NAT gateway |

This architecture diagram depicts what you create in the lab

< Previous

Next >

This architecture diagram depicts what you create in the lab.

52

Student Guide


awsacademy.instructure.com/courses/2401/modules/items/278437

ACFv2EN-2401 > Modules > Module 5 - Networking and Content Delivery > Student Guide

Home
Announcements
Modules
Discussions
Grades

Account
Dashboard
Courses
Calendar
Inbox
History
Help

100-ACCLFO-20-EN-M05SG ~ 30 minutes 53 of 79 Close



It is now time to start the lab. It should take you approximately 30 minutes to complete the lab.

Previous Next

It is now time to start the lab. It should take you approximately 30 minutes to complete the lab.

Student Guide

awsacademy.instructure.com/courses/2401/modules/items/278437

aws ACFv2EN-2401 > Modules > Module 5 - Networking and Content Delivery > Student Guide

100-ACCLFO-20-EN-M05SG 54 of 79

Home
Announcements
Modules
Discussions
Grades

Account
Dashboard
Courses
Calendar
Inbox
History
Help

aws academy

Lab debrief:
Key takeaways

© 2019 Amazon Web Services, Inc. or its affiliates. All rights reserved.

In this lab, you:

Previous Next

In this lab, you:

- Created an Amazon VPC.
- Created additional subnets.
- Created an Amazon VPC security group.
- Launched a web server instance on Amazon EC2.

Student Guide

awsacademy.instructure.com/courses/2401/modules/items/278437

aws

Account

Dashboard

Courses

Calendar

Inbox

History

Help

ACFv2EN-2401 > Modules > Module 5 - Networking and Content Delivery > Student Guide

Home

Announcements

Modules

Discussions

Grades

100-ACCLFO-20-EN-M05SG

< > 55 of 79

+ - <->

Close

Module 5: Networking and Content Delivery

Section 5: Amazon Route 53

© 2019 Amazon Web Services, Inc. or its affiliates. All rights reserved.

aws academy

Section 5: Amazon Route 53

< Previous

Next >

Student Guide

awsacademy.instructure.com/courses/2401/modules/items/278437

ACFv2EN-2401 > Modules > Module 5 - Networking and Content Delivery > Student Guide

aws

Account

Dashboard

Courses

Calendar

Inbox

History

Help

Home


Announcements

Modules

Discussions

Grades

100-ACCLFO-20-EN-M05SG < > 58 of 79 + - <-> Close



Amazon Route 53

- Is a highly available and scalable Domain Name System (DNS) web service
- Is used to route end users to internet applications by translating names (like [www.example.com](#)) into numeric IP addresses (like 192.0.2.1) that computers use to connect to each other
- Is fully compliant with IPv4 and IPv6
- Connects user requests to infrastructure running in AWS and also outside of AWS
- Is used to check the health of your resources
- Features traffic flow
- Enables you to register domain names

Amazon Route 53 is a highly available and scalable cloud [Domain Name System \(DNS\)](#) web

Previous

Next

Amazon Route 53 is a highly available and scalable cloud

Domain Name System (DNS) web

service. It is designed to give developers and busi

nesses a reliable and cost-effective way to

route users to internet applications by translating

names (like `www.example.com`) into the

numeric IP addresses (like `192.0.2.1`) that computers use to connect to

each other. In addition, Amazon Route 53 is fully compliant with IPv6.

Amazon Route 53 effectively connects user requests

to infrastructure running in AWS—such

as Amazon EC2 instances, Elastic Load Balancing load balancers, or

Amazon S3 buckets—and

can also be used to route users to infrastructure t

hat is outside of AWS.

You can use Amazon Route 53 to configure DNS health

checks so you that can route traffic to

healthy endpoints or independently monitor the health of your application

and its endpoints.

Amazon Route 53 traffic flow helps you manage traffic globally through

several routing types,

which can be combined with DNS failover to enable various low-latency,

fault-tolerant

architectures. You can use Amazon Route 53 traffic

flow's simple visual editor to manage

how your users are routed to your application's end

points—whether in a single AWS Region

or distributed around the globe.

Amazon Route 53 also offers Domain Name Registration—you can

purchase and manage

domain names (like `example.com`), and Amazon Route 53 will

automatically configure

DNS settings for your domains.

Student Guide

awsacademy.instructure.com/courses/2401/modules/items/278437

ACFv2EN-2401 > Modules > Module 5 - Networking and Content Delivery > Student Guide

100-ACCLFO-20-EN-M05SG < > 57 of 79 + - <-> Close

Home
Announcements
Modules
Discussions
Grades

Account
Dashboard
Courses
Calendar
Inbox
History
Help

Amazon Route 53 DNS resolution

```

graph LR
    User((User)) -- "Requests www.example.com" --> DNSResolver[DNS resolver]
    DNSResolver -- "Checks with Route 53 for IP address" --> Route53[Amazon Route 53]
    Route53 -- "Returns IP address 192.0.2.0" --> DNSResolver
    DNSResolver -- "Returns IP address 192.0.2.0" --> User
  
```

© 2019 Amazon Web Services, Inc. or its affiliates. All rights reserved.

Previous Next

Here is the basic pattern that Amazon Route 53 follows when a user initiates a DNS request. The DNS resolver checks with your domain in Route 53, gets the IP address, and returns it to the user.

57

Student Guide

awsacademy.instructure.com/courses/2401/modules/items/278437

ACFv2EN-2401 > Modules > Module 5 - Networking and Content Delivery > Student Guide

100-ACCLFO-20-EN-M05SG < > 58 of 79 + - <-> Close

Home
Announcements
Modules
Discussions
Grades

Account
Dashboard
Courses
Calendar
Inbox
History
Help

Amazon Route 53 supported routing

- **Simple routing** – Use in single-server environments
- **Weighted round robin routing** – Assign weights to resource record sets to specify the frequency
- **Latency routing** – Help improve your global applications
- **Geolocation routing** – Route traffic based on location of your users
- **Geoproximity routing** – Route traffic based on location of your resources
- **Failover routing** – Fail over to a backup site if your primary site becomes unreachable
- **Multivalue answer routing** – Respond to DNS queries with up to eight healthy records selected at random

© 2019 Amazon Web Services, Inc. or its affiliates. All rights reserved.

Previous Next

Amazon Route 53 supports several types of routing policies, which determine how Amazon Route 53 responds to queries:

-

Simple routing (round robin)

- Use for a single resource that performs a given function for your domain (such as a web server that serves content for the example.com website).

-

Weighted round robin routing –

Use to route traffic to multiple resources in proportions that you specify. Enables you to assign weights to resource record sets to specify the frequency with which different responses are served

. You might want to use this

capability to do A/B testing, which is when you send a small portion of traffic to a server

where you made a software change. For instance, suppose you have two record sets that

are associated with one DNS name: one with weight 3

and one with weight 1. In this case,

75 percent of the time, Amazon Route 53 will return

the record set with weight 3, and 25

percent of the time, Amazon Route 53 will return the record set with weight 1. Weights

can be any number between 0 and 255.

-

Latency routing (LBR) –

Use when you have resources in multiple AWS Regions and you want to route traffic to the Region that provides the best latency. Latency routing works

by routing your customers to the AWS endpoint (for example, Amazon EC2 instances,

Elastic IP addresses, or load balancers) that provides the fastest experience based on

actual performance measurements of the different AW

S Regions where your application

runs.

-

Geolocation routing –

Use when you want to route traffic based on the location of your users. When you use geolocation routing, you can localize your content and present some or all of your website in the language of your users. You can also use geolocation routing to restrict the distribution of content to only the locations where you have distribution rights. Another possible use is for balancing the load across endpoints in a predictable, easy-to-manage way, so that each user location is consistently routed to the same endpoint.

-

Geoproximity routing

– Use when you want to route traffic based on the location of your resources and, optionally, shift traffic from resources in one location to resources in another.

-

Failover routing (DNS failover) –

Use when you want to configure active-passive failover.

Amazon Route 53 can help detect an outage of your website and redirect your users to alternate locations where your application is operating properly. When you enable this feature, Amazon Route 53 health-checking agents will monitor each location or endpoint of your application to determine its availability.

You can take advantage of this feature to increase the availability of your customer-facing application.

-

Multivalue answer routing– Use when you want Route 53 to respond to DNS queries with up to eight healthy records that are selected at random. You can configure Amazon Route 53 to return multiple values—such as IP addresses for your web servers—in response to DNS queries. You can specify multiple values for almost any record, but multivalue answer routing also enables you to check the health of each resource so that Route 53 returns only values for healthy resources. It's not a substitute for a load balancer, but the ability to return multiple health-checkable IP addresses is a way to use DNS to improve availability and load balancing.

Student Guide

awsacademy.instructure.com/courses/2401/modules/items/278437

ACFv2EN-2401 > Modules > Module 5 - Networking and Content Delivery > Student Guide

aws

Account

Dashboard

Courses

Calendar

Inbox

History

Help

Home

Announcements

Modules

Discussions

Grades

100-ACCLFO-20-EN-M05SG

< > 60 of 79

+ - <->

Close

Use case: Multi-region deployment

Amazon Route 53

some-elb-name.us-west-2.elb.amazonaws.com

User

some-elb-name.ap-southeast-2.elb.amazonaws.com

| Name | Type | Value |
|-------------|-------|--|
| example.com | ALIAS | some-elb-name.us-west-2.elb.amazonaws.com |
| example.com | ALIAS | some-elb-name.ap-southeast-2.elb.amazonaws.com |

Multi-Region deployment is an example use case for Amazon Route 53. With Amazon Route

< Previous

Next >

Multi-Region deployment is an example use case for Amazon Route 53. With Amazon Route 53, the user is automatically directed to the Elastic Load Balancing load balancer that's closest to the user.

The benefits of multi-region deployment of Route 53 include:

-

Latency-based routing to the Region

-

Load balancing routing to the Availability Zone

The screenshot shows a web browser window with the URL `awsacademy.instructure.com/courses/2401/modules/items/278437`. The page is titled "Student Guide" and shows a navigation sidebar on the left with links to Home, Account, Dashboard, Courses, Calendar, Inbox, History, and Help. The main content area displays a slide titled "Amazon Route 53 DNS failover" with the AWS Academy logo. The slide text reads: "Improve the availability of your applications that run on AWS by:" followed by a bulleted list: "• Configuring backup and failover scenarios for your own applications", "• Enabling highly available multi-region architectures on AWS", and "• Creating health checks". Below the text is a screenshot of the "Advanced configuration" dialog box for Amazon Route 53, showing options for "Request interval" (Standard (30 seconds) selected), "Failure threshold" (5), "Bring matching" (No selected), "Latency graphs" (checked), and "Invert health check status" (unchecked). The slide is part of a presentation, with a "100-ACCLFO-20-EN-M05SG" identifier and a "61 of 79" page indicator.

Amazon Route 53 enables you to improve the availability of your applications that run on AWS by:

-
- Configuring backup and failover scenarios for your own applications.
-
- Enabling highly available multi-Region architectures on AWS.
-
- Creating health checks to monitor the health and performance of your web applications, web servers, and other resources. Each health check that you create can monitor one of the following—the health of a specified resource, such as a web server; the status of other health checks; and the status of an Amazon CloudWatch alarm.

Student Guide

awsacademy.instructure.com/courses/2401/modules/items/278437

ACFv2EN-2401 > Modules > Module 5 - Networking and Content Delivery > Student Guide

Account

Dashboard

Courses

Calendar

Inbox

History

Help

Home

Announcements

Modules

Discussions

Grades

100-ACCLFO-20-EN-M05SG

< > 62 of 79

Close

DNS failover for a multi-tiered web application

Record Sets

- CNAME www
- elastic_load_balancing Routing Policy = Failover
- Record Type = Primary
- Amazon S3 website Routing Policy = Failover
- Record Type = Secondary

User

Amazon Route 53

Primary

Secondary

Amazon S3 static website

AWS Cloud

Availability Zone A

Availability Zone B

Auto Scaling group

Amazon EC2

Amazon RDS

Amazon Relational Database Service (Amazon RDS) instance

Here, you see how DNS failover works in a typical architecture for a multi-tiered web.

< Previous

Next >

Here, you see how DNS failover works in a typical architecture for a multi-tiered web application. Route 53 passes traffic to a load balancer, which then distributes traffic to a fleet of EC2 instances.

You can do the following tasks with Route 53 to ensure high availability:

1. Create two DNS records for the Canonical Name Record (CNAME) `www`

with a routing

policy of

Failover Routing

- . The first record is the primary route policy, which points to the load balancer for your web application. The second record is the secondary route policy, which points to your static Amazon S3 website.

2. Use Route 53 health checks to make sure that the primary is running. If it is, all traffic

defaults to your web application stack. Failover to

the static backup site would be

triggered if either the web server goes down (or stops responding), or the database

instance goes down.

Student Guide

awsacademy.instructure.com/courses/2401/modules/items/278437

aws

Account

Dashboard

Courses

Calendar

Inbox

History

Help

ACFv2EN-2401 > Modules > Module 5 - Networking and Content Delivery > Student Guide

Home

Announcements

Modules

Discussions

Grades


100-ACCLFO-20-EN-M05SG

< > 63 of 79

+ - <->

Close

Section 5 key takeaways



aws academy

- Amazon Route 53 is a highly available and scalable cloud DNS web service that translates domain names into numeric IP addresses.
- Amazon Route 53 supports several types of routing policies.
- Multi-Region deployment improves your application's performance for a global audience.
- You can use Amazon Route 53 failover to improve the availability of your applications.

Some key takeaways from this section of the module include:

- Amazon Route 53 is a highly available and scalable cloud DNS web service that translates

< Previous

Next >

Some key takeaways from this section of the module include:

-

Amazon Route 53 is a highly available and scalable cloud DNS web service that translates domain names into numeric IP addresses.

-

Amazon Route 53 supports several types of routing policies.

-

Multi-Region deployment improves your application's performance for a global audience.

-

You can use Amazon Route 53 failover to improve the availability of your applications.

Student Guide

awsacademy.instructure.com/courses/2401/modules/items/278437

ACFv2EN-2401 > Modules > Module 5 - Networking and Content Delivery > Student Guide

aws

Account

Dashboard

Courses

Calendar

Inbox

History

Help

Home

Announcements

Modules

Discussions

Grades

100-ACCLFO-20-EN-M05SG < > 64 of 79 + - <-> Close

Module 5: Networking and Content Delivery

Section 6: Amazon CloudFront

© 2019 Amazon Web Services, Inc. or its affiliates. All rights reserved. aws academy

Section 6: Amazon CloudFront

The purpose of networking is to share information between connected resources. So far in

< Previous

Next >

Section 6: Amazon CloudFront

The purpose of networking is to share information between connected resources. So far in this module, you learned about VPC networking with Amazon VPC. You learned about the different options for connecting your VPC to the internet, to remote networks, to other VPCs, and to AWS services.

Content delivery occurs over networks, too—for example, when you stream a movie from your favorite streaming service. In this final section, you learn about Amazon CloudFront, which is a content delivery network (CDN) service.

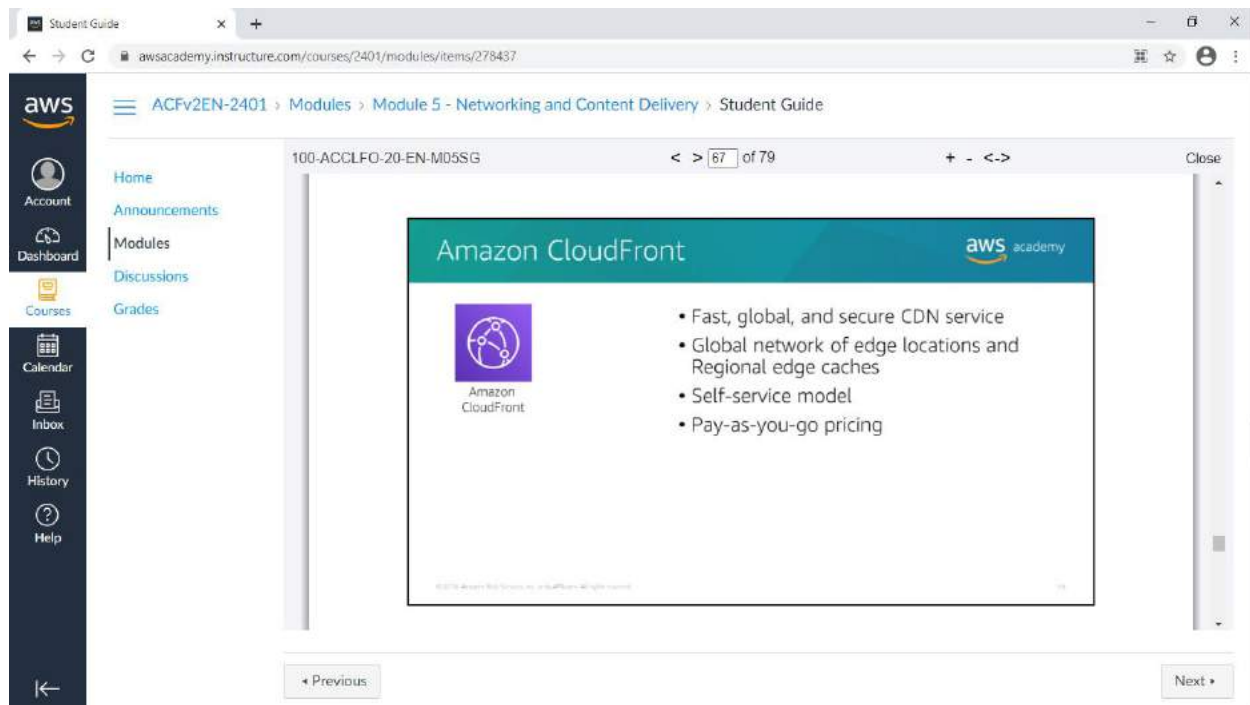
As explained earlier in this module when you were learning about AWS Direct Connect, one of the challenges of network communication is network performance. When you browse a website or stream a video, your request is routed through many different networks to reach an origin server. The origin server (or origin) stores the original, definitive versions of the objects (webpages, images, and media files). The number of network hops and the distance that the request must travel significantly affect the performance and responsiveness of the website. Further, network latency is different in various geographic locations. For these reasons, a content delivery network might be the solution.

The screenshot shows a web browser window with the URL `awsacademy.instructure.com/courses/2401/modules/items/278437`. The page is titled "Student Guide" and displays a navigation sidebar on the left with links to Home, Account, Dashboard, Courses, Calendar, Inbox, History, and Help. The main content area shows a slide titled "Content delivery network (CDN)" with the AWS Academy logo. The slide lists the following points:

- Is a globally distributed system of caching servers
- Caches copies of commonly requested files (static content)
- Delivers a local copy of the requested content from a nearby cache edge or Point of Presence
- Accelerates delivery of dynamic content
- Improves application performance and scaling

At the bottom of the slide, it says "© 2019 Amazon Web Services, Inc. or its affiliates. All rights reserved." The slide is part of a presentation, with navigation controls for "Previous" and "Next" visible at the bottom.

A content delivery network (CDN) is a globally distributed system of caching servers. A CDN caches copies of commonly requested files (static content, such as Hypertext Markup Language, or HTML; Cascading Style Sheets, or CSS; JavaScript; and image files) that are hosted on the application origin server. The CDN delivers a local copy of the requested content from a cache edge or Point of Presence that provides the fastest delivery to the requester. CDNs also deliver dynamic content that is unique to the requester and is not cacheable. Having a CDN deliver dynamic content improves application performance and scaling. The CDN establishes and maintains secure connections closer to the requester. If the CDN is on the same network as the origin, routing back to the origin to retrieve dynamic content is accelerated. In addition, content such as form data, images, and text can be ingested and sent back to the origin, thus taking advantage of the low-latency connections and proxy behavior of the PoP.



Amazon CloudFront is a fast CDN service that securely delivers data, videos, applications, and application programming interfaces (APIs) to customers globally with low latency and high transfer speeds. It also provides a developer-friendly environment. Amazon CloudFront delivers files to users over a global network of edge locations and Regional edge caches. Amazon CloudFront is different from traditional content delivery solutions because it enables you to quickly obtain the benefits of high-performance content delivery without negotiated contracts, high prices, or minimum fees. Like other AWS services, Amazon CloudFront is a self-service offering with pay-as-you-go pricing.

Student Guide

awsacademy.instructure.com/courses/2401/modules/items/278437

100-ACCLFO-20-EN-M05SG68 of 79Close

aws

Account

Dashboard

Courses

Calendar

Inbox

History

Help

Home

Announcements

Modules

Discussions

Grades

Amazon CloudFront infrastructure

Edge locations

Multiple edge locations

Regional edge caches

• Edge locations – Network of data centers that CloudFront uses to serve popular content quickly to customers.

• Regional edge cache – CloudFront location that caches content that is not popular enough to stay at an edge location. It is located between the origin server and the global edge location.

© 2019 Amazon Web Services, Inc. or its affiliates. All rights reserved.

Previous

Next

Student Guide

awsacademy.instructure.com/courses/2401/modules/items/278437

100-ACCLFO-20-EN-M05SG68 of 79Close

aws

AccountDashboardCoursesCalendarInboxHistoryHelp

ACFv2EN-2401ModulesModule 5 - Networking and Content DeliveryStudent Guide


HomeAnnouncementsModulesDiscussionsGrades

Amazon CloudFront infrastructure

Edge locationsMultiple edge locationsRegional edge caches

Edge locations – Network of data centers that CloudFront uses to serve popular content quickly to customers.

Regional edge cache – CloudFront location that caches content that is not popular enough to stay at an edge location. It is located between the origin server and the global edge location.



Previous

Next

Amazon CloudFront delivers content through a worldwide network of data centers that are called edge locations

. When a user requests content that you serve with CloudFront, the user is routed to the edge location that provides the lowest latency (or time delay) so that content

is delivered with the best possible performance. CloudFront edge locations are designed to serve popular content quickly to your viewers.

As objects become less popular, individual edge locations might remove those objects to

make room for more popular content. For the less popular content, CloudFront has

Regional edge caches.

Regional edge caches are CloudFront locations that are deployed globally and

are close to your viewers. They are located between your origin server and the global edge locations that serve content directly to viewers. A

Regional edge cache has a larger cache than an individual edge location, so objects remain in the Regional edge cache longer. More

of your content remains closer to your viewers, which reduces the need for CloudFront to go

back to your origin server and improves overall performance for viewers.

For more information about how Amazon CloudFront works, see

How CloudFront Delivers

Contents in the AWS Documentation.

Student Guide

awsacademy.instructure.com/courses/2401/modules/items/278437

ACFv2EN-2401 > Modules > Module 5 - Networking and Content Delivery > Student Guide

Account

Dashboard

Courses

Calendar

Inbox

History

Help

Home

Announcements

Modules

Discussions

Grades

100-ACCLFO-20-EN-M05SG

< > 69 of 79

+ - <->

Close

Amazon CloudFront benefits

- Fast and global
- Security at the edge
- Highly programmable
- Deeply integrated with AWS
- Cost-effective

© 2019 Amazon Web Services, Inc. or its affiliates. All rights reserved.

Amazon CloudFront provides the following benefits:

- Fast and global: Amazon CloudFront is managed, scaled, and globally distributed. To

< Previous

Next >

Amazon CloudFront provides the following benefits:

-

Fast and global

- Amazon CloudFront is massively scaled and globally distributed. To

deliver content to end users with low latency, Amazon CloudFront uses a global network

that consists of edge locations and regional caches

-

-

Security at the edge

- Amazon CloudFront provides both network-level and application-level protection. Your traffic and applications benefit through various built-in protections,

such as AWS Shield Standard, at no additional cost.

You can also use configurable features,

such as AWS Certificate Manager (ACM), to create and manage custom Secure Sockets Layer (SSL) certificates at no extra cost.

-

Highly programmable

- Amazon CloudFront features can be customized for specific application requirements. It integrates with Lambda@Edge so that you can run custom code across AWS locations worldwide, which enables you to move complex application logic closer to users to improve responsiveness. The CDN also supports integrations with other tools and automation interfaces for DevOps. It offers continuous integration and

continuous delivery (CI/CD) environments.

-

Deeply integrated with AWS

- Amazon CloudFront is integrated with AWS, with both physical locations that are directly connected to the AWS Global Infrastructure and other AWS services. You can use APIs or the AWS Management Console to programmatically configure all features in the CDN.

69

-

Cost-effective

– Amazon CloudFront is cost-effective because it has no minimum commitments and charges you only for what you use. Compared to self-hosting, Amazon CloudFront avoids the expense and complexity of operating a network of cache servers in multiple sites across the internet. It eliminates the need to overprovision capacity to serve potential spikes in traffic. Amazon CloudFront also uses techniques like collapsing simultaneous viewer requests at an edge location for the same file into a single request to your origin server. The result is reduced load on your origin servers and reduced need to scale your origin infrastructure, which can result in further cost savings. If you use AWS origins such as Amazon Simple Storage Service (Amazon S3) or Elastic Load Balancing, you pay only for storage costs, not for any data transferred between these services and CloudFront.

Student Guide

awsacademy.instructure.com/courses/2401/modules/items/278437

ACFv2EN-2401 > Modules > Module 5 - Networking and Content Delivery > Student Guide

aws

Account

Dashboard

Courses

Calendar

Inbox

History

Help

Home

Announcements

Modules

Discussions

Grades

100-ACCLFO-20-EN-M05SG < > 71 of 79 + - <-> Close

Amazon CloudFront pricing

Data transfer out

- Charged for the volume of data transferred out from Amazon CloudFront edge location to the internet or to your origin.

HTTP(S) requests

- Charged for number of HTTP(S) requests.

Invalidation requests

- No additional charge for the first 1,000 paths that are requested for invalidation each month. Thereafter, \$0.005 per path that is requested for invalidation.

Dedicated IP custom SSL

- \$600 per month for each custom SSL certificate that is associated with one or more CloudFront distributions that use the Dedicated IP version of custom SSL certificate support.

© 2019 Amazon Web Services, Inc. or its affiliates. All rights reserved.

Amazon CloudFront charges are based on actual usage of the service in four areas:

< Previous

Next >

Amazon CloudFront charges are based on actual usage of the service in four areas:

-

Data transfer out –You are charged for the volume of data that is transferred out from

Amazon CloudFront edge locations, measured in GB, to the internet or to your origin (both AWS origins and other origin servers). Data transfer usage is totaled separately for specific geographic regions, and then cost is calculated based on pricing tiers for each area. If you use other AWS services as the origins of your files

, you are charged separately for your use of those services, including storage and compute hours.

-

HTTP(S) requests –

You are charged for the number of HTTP(S) requests that are made to

Amazon CloudFront for your content.

-

Invalidation requests –

You are charged per path in your invalidation request. A path that is listed in your invalidation request represents the URL (or multiple URLs if the path contains a wildcard character) of the object that you want to invalidate from CloudFront

cache. You can request up to 1,000 paths each month

from Amazon CloudFront at no additional charge. Beyond the first 1,000 paths, you are charged per path that is listed in your invalidation requests.

-

Dedicated IP custom Secure Sockets Layer (SSL) –

You pay \$600 per month for each custom SSL certificate that is associated with one or more CloudFront distributions that

use the Dedicated IP version of custom SSL certificate support. This monthly fee is prorated by the hour. For example, if your custom SSL certificate was associated with at

least one CloudFront distribution for just 24 hours

(that is, 1 day) in the month of June,

your total charge for using the custom SSL certificate feature in June is $(1 \text{ day} / 30 \text{ days}) * \$600 = \$20$.

For the latest pricing information, see the
Amazon CloudFront pricing page

.

The screenshot shows the AWS Academy Student Guide interface. The left sidebar contains navigation links: Home, Announcements, Modules, Discussions, Grades, Account, Dashboard, Courses, Calendar, Inbox, History, and Help. The main content area displays a slide titled 'Section 6 key takeaways' with a key icon and the word 'Takeaway'. The slide content includes the following bullet points:

- A CDN is a globally distributed system of caching servers that accelerates delivery of content.
- Amazon CloudFront is a fast CDN service that securely delivers data, videos, applications, and APIs over a global infrastructure with low latency and high transfer speeds.
- Amazon CloudFront offers many benefits.

The slide is part of a presentation titled '100-ACCLFO-20-EN-M05SG' and is slide 73 of 79. Navigation buttons for 'Previous' and 'Next' are visible at the bottom.

Some key takeaways from this section of the module include:

-

A CDN is a globally distributed system of caching servers that accelerates delivery of content.

-

Amazon CloudFront is a fast CDN service that securely delivers data, videos, applications, and APIs over a global infrastructure with low latency and high transfer speeds.

-

Amazon CloudFront offers many benefits, including:

-

Fast and global

-

Security at the edge

-

Highly programmable

-

Deeply integrated with AWS

-

Cost-effective

Student Guide

awsacademy.instructure.com/courses/2401/modules/items/278437

ACFv2EN-2401 > Modules > Module 5 - Networking and Content Delivery > Student Guide

100-ACCLFO-20-EN-M05SG < > 74 of 79 + - <-> Close

Module 5: Networking and Content Delivery

Module wrap-up

aws academy

© 2019 Amazon Web Services, Inc. or its affiliates. All rights reserved.

It's now time to review the module, and wrap up with a knowledge check and a discussion of a practice certification exam question.

◀ Previous Next ▶

It's now time to review the module, and wrap up with a knowledge check and a discussion of a practice certification exam question.

Student Guide

awsacademy.instructure.com/courses/2401/modules/items/278437

aws

Account

Dashboard

Courses

Calendar

Inbox

History

Help

ACFv2EN-2401 > Modules > Module 5 - Networking and Content Delivery > Student Guide

Home

Announcements

Modules

Discussions

Grades

100-ACCLFO-20-EN-M05SG

< > 75 of 79

+ - <->

Close

Module summary

aws academy

In summary, in this module you learned how to:

- Recognize the basics of networking
- Describe virtual networking in the cloud with Amazon VPC
- Label a network diagram
- Design a basic VPC architecture
- Indicate the steps to build a VPC
- Identify security groups
- Create your own VPC and added additional components to it to produce a customized network
- Identify the fundamentals of Amazon Route 53
- Recognize the benefits of Amazon CloudFront

In summary, in this module you learned how to:

- Recognize the basics of networking

< Previous

Next >

In summary, in this module you learned how to:

-

Recognize the basics of networking

-

Describe virtual networking in the cloud with Amazon VPC

-

Label a network diagram

-

Design a basic VPC architecture

-

Indicate the steps to build a VPC

-

Identify security groups

-

Create your own VPC and added additional components to it to produce a customized network

-

Identify the fundamentals of Amazon Route 53

-

Recognize the benefits of Amazon CloudFront

Student Guide

awsacademy.instructure.com/courses/2401/modules/items/278437

aws

Account

Dashboard

Courses

Calendar

Inbox

History

Help

ACFv2EN-2401 > Modules > Module 5 - Networking and Content Delivery > Student Guide

Home

Announcements

Modules

Discussions

Grades

100-ACCLFO-20-EN-M05SG

< > 77 of 79

+ - <->

Close

Sample exam question

aws academy

Which AWS networking service enables a company to create a virtual network within AWS?

A. AWS Config

B. Amazon Route 53

C. AWS Direct Connect

D. Amazon VPC

© 2019 Amazon Web Services, Inc. or its affiliates. All rights reserved.

Look at the answer choices and rule them out based on the keywords that were previously

< Previous

Next >

The screenshot shows the AWS Academy Student Guide interface. The browser address bar displays the URL `awsacademy.instructure.com/courses/2401/modules/items/278437`. The page header includes the AWS logo and navigation links: `ACFv2EN-2401 > Modules > Module 5 - Networking and Content Delivery > Student Guide`. A sidebar on the left contains navigation options: Home, Announcements, Modules, Discussions, Grades, Calendar, Inbox, History, and Help. The main content area shows a slide titled "Additional resources" with the AWS Academy logo. The slide lists five resources:

- [Amazon VPC overview page](#)
- [Amazon Virtual Private Cloud Connectivity Options](#) whitepaper
- [One to Many: Evolving VPC Design](#) AWS Architecture blog post
- [Amazon VPC User Guide](#)
- [Amazon CloudFront overview page](#)

At the bottom of the slide, there is a small copyright notice: "© 2020 Amazon Web Services, Inc. or its affiliates. All rights reserved." Below the slide, there are navigation buttons: "Previous" and "Next".

If you want to learn more about the topics covered in this module, you might find the following additional resources helpful:

- Amazon VPC overview page
- Amazon Virtual Private Cloud Connectivity Options whitepaper
- One to Many: Evolving VPC Design AWS Architecture blog post
- Amazon VPC User Guide
- Amazon CloudFront overview page